16-824:Visual Learning and Recognition

Many slides from A. Farhadi, A. Efros

Course Information

- Time:
  - Monday, Wednesday 12:00-1:20
- Location:
  - Gates 4307
- Office Hours:
  - Email me for appointments
- Contact:
  - abhinavg@cs , EDSH 213
- Website:
  -

People - Instructor

- Abhinav Gupta
- Ph.D. 2009, University of Maryland

Pitcher pitches the ball before Batter hits. Batter hits and then simultaneously Batter runs to base and Fielder runs towards the ball. Fielder runs towards the ball and then Fielder catches the ball. Fielder catches the ball and then Fielder throws to the base. Fielder at Base catches the ball at base after Fielder throws to the base.

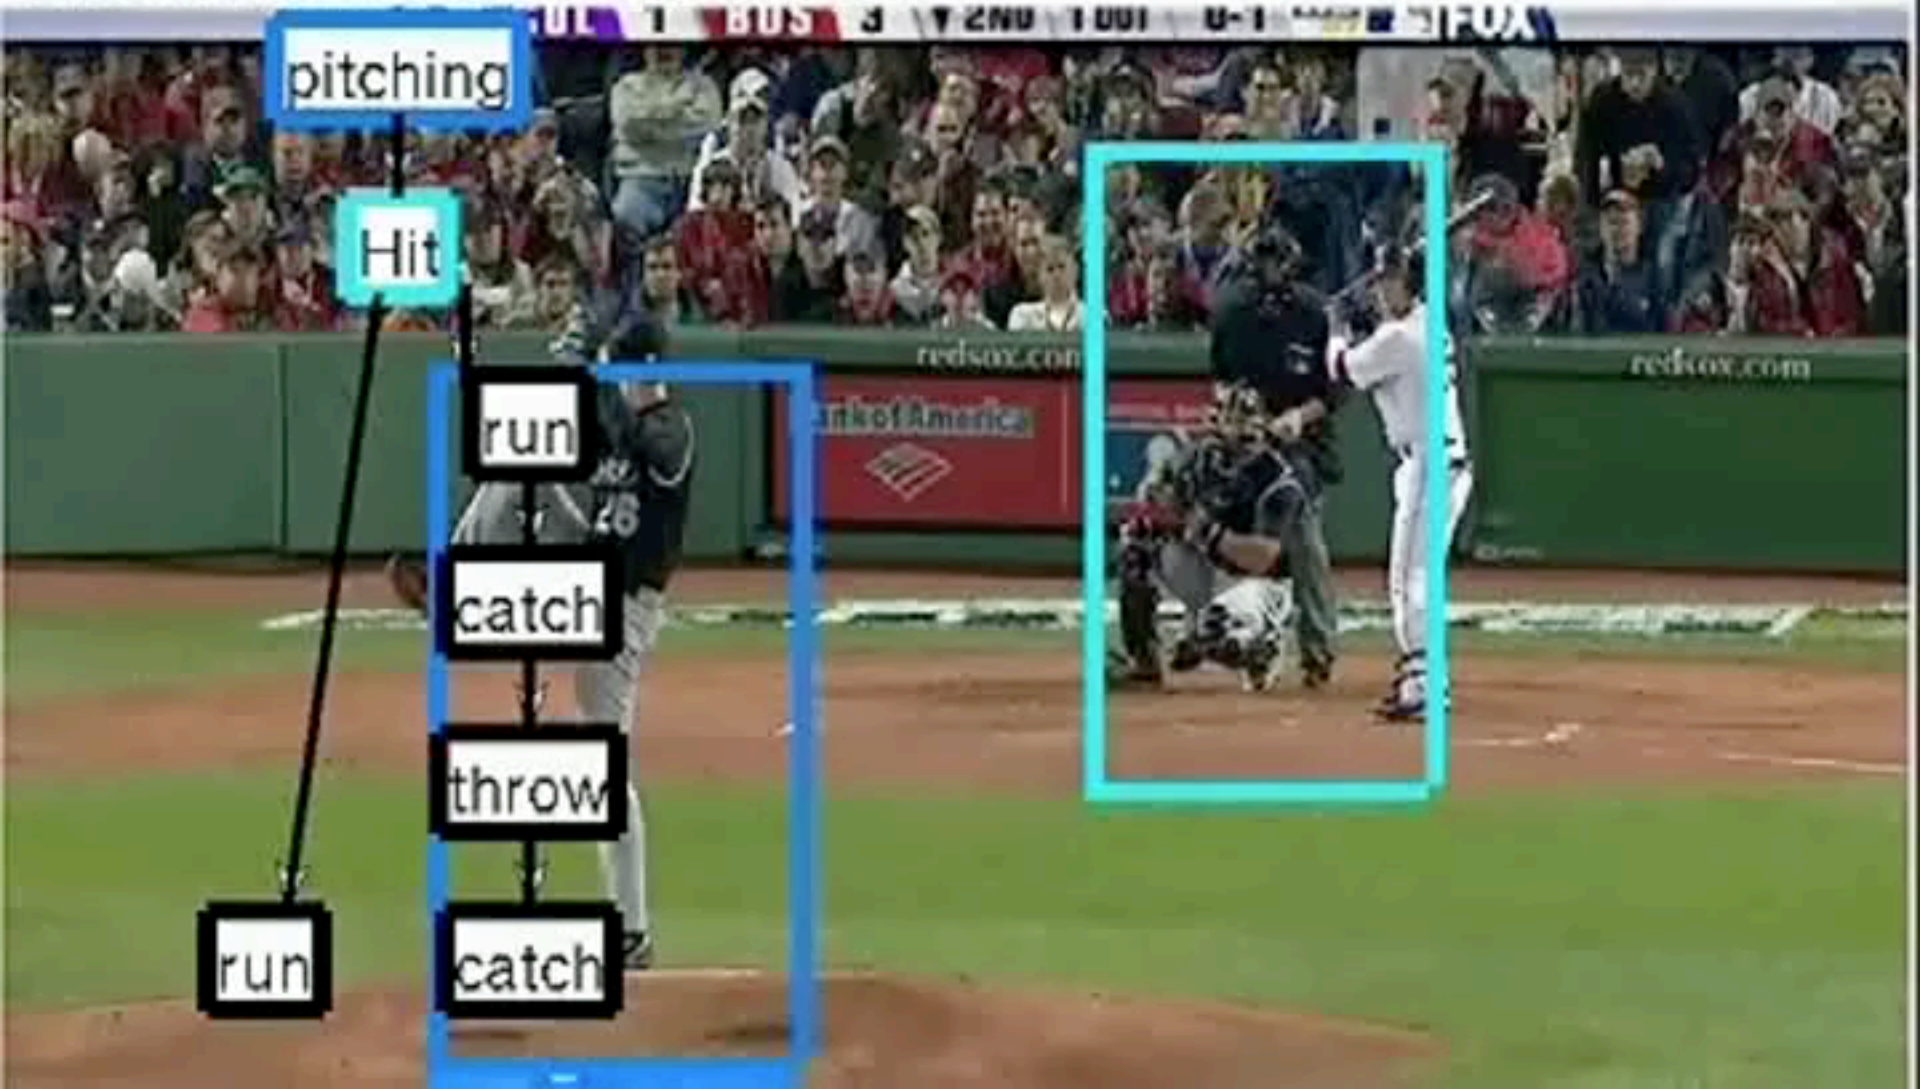

Image captioning — today

a man riding a motorcycle on a city street

a baseball glove with a ball in it

People

- Abhinav Gupta
- Ph.D. 2009, University of Maryland
- Postdoctoral Fellow, Carnegie Mellon University, 2009-11

People

- Abhinav Gupta
- Ph.D. 2009, University of Maryland
- Postdoctoral Fellow, Carnegie Mellon University, 2009-11

Recently

- Lifelong Learning - NEIL

NEIL


Build Visual Models
Learn Common-Sense Relationships
Bean
Falcon
Wheel is a part of Car
Eye is a part of Baby
Helicopter is found in Airfield
Zebra is found in Savanna
Trading Floor are Crowded
Wheels have round shape

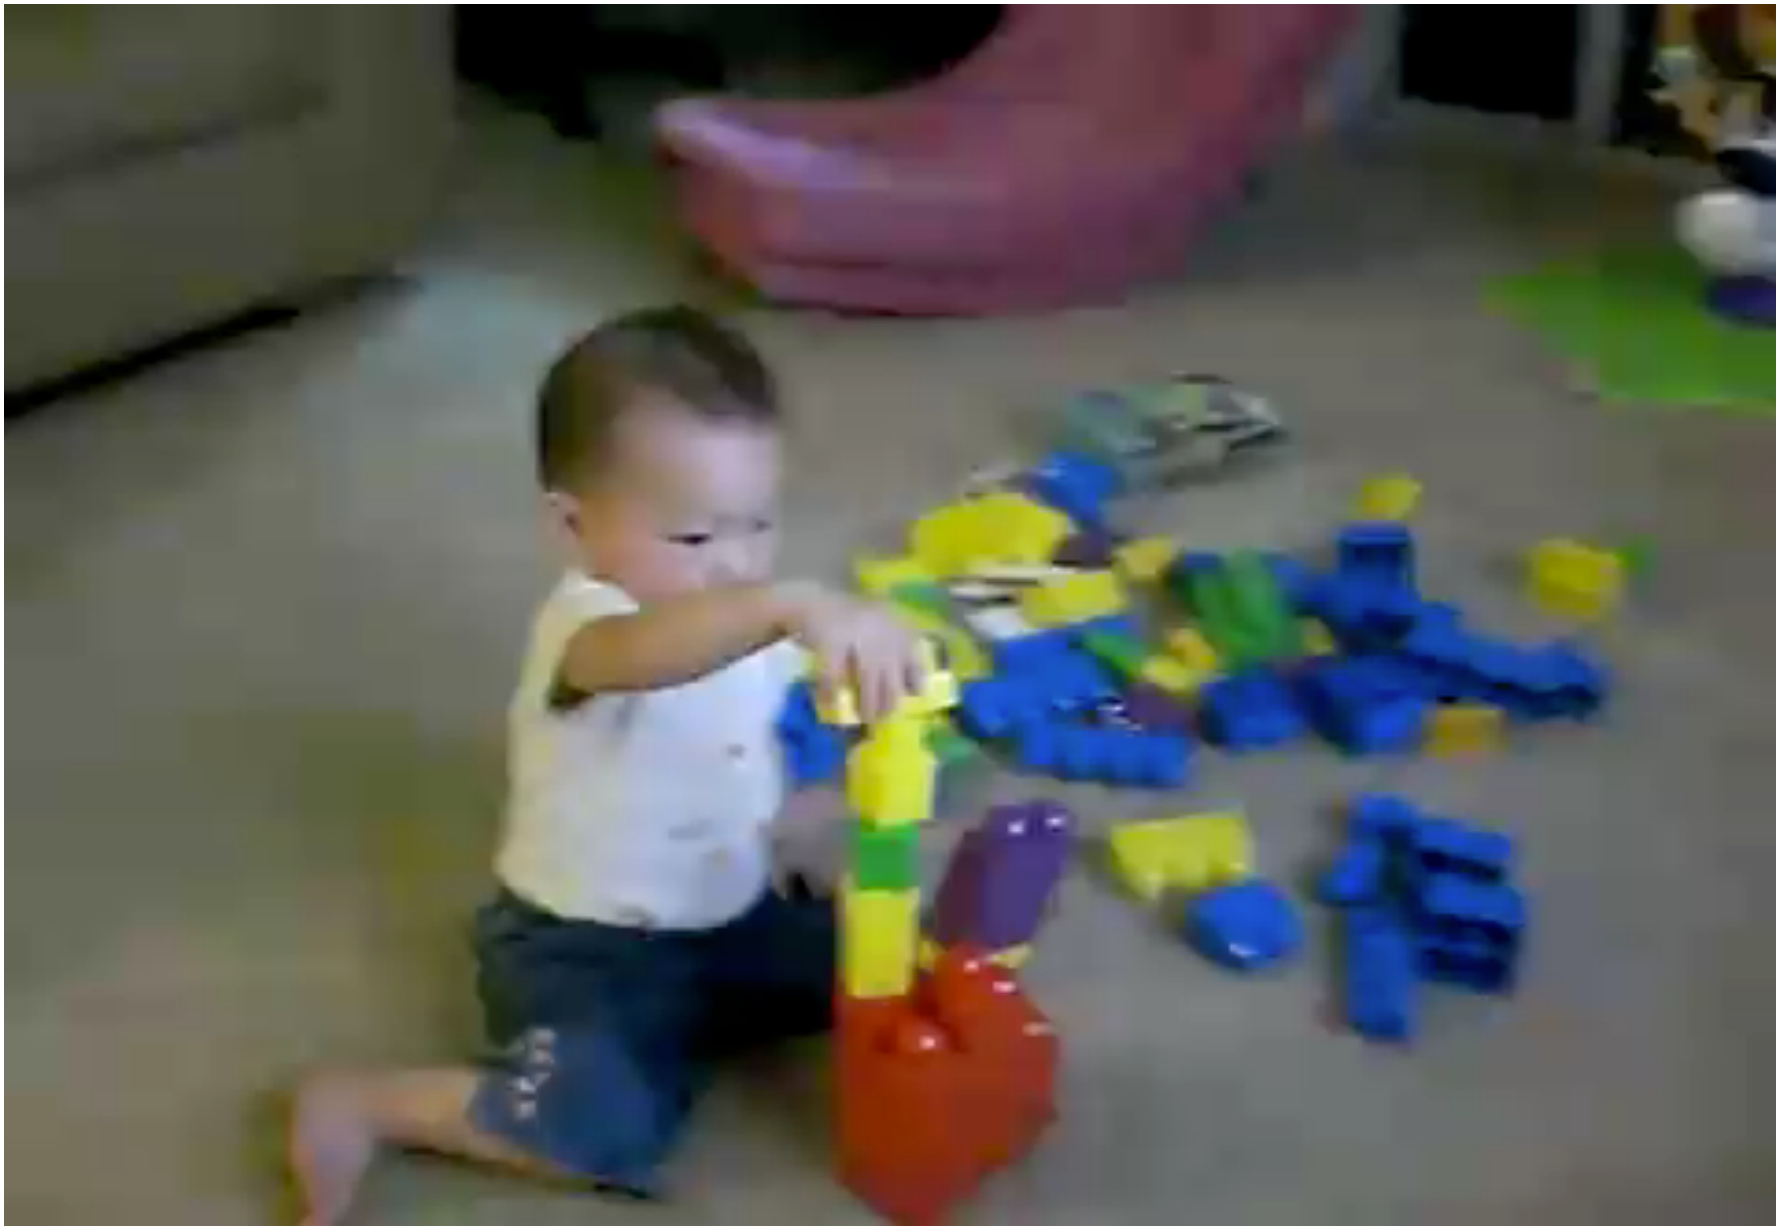

People

- Abhinav Gupta
- Ph.D. 2009, University of Maryland
- Postdoctoral Fellow, Carnegie Mellon University, 2009-11

Recently

- Lifelong Learning - NEIL
- Learning from Physical Interactions

Teaching Assistants

Rohit Girdhar

Second year PhD student

Research Interests:

Action Recognition, Pose Estimation

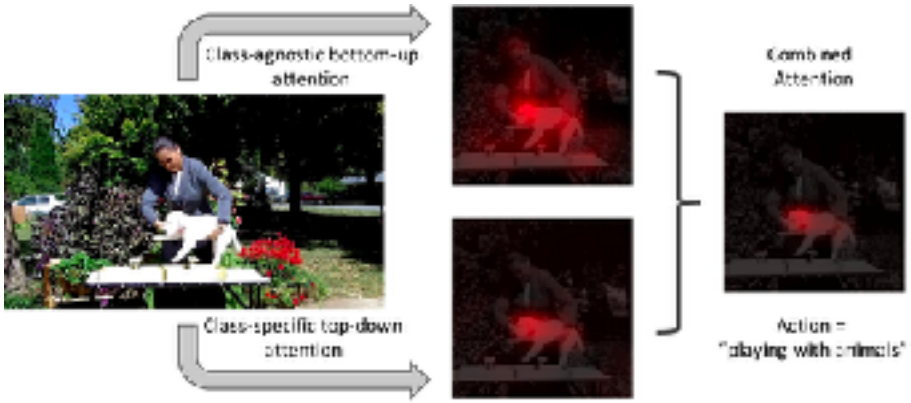

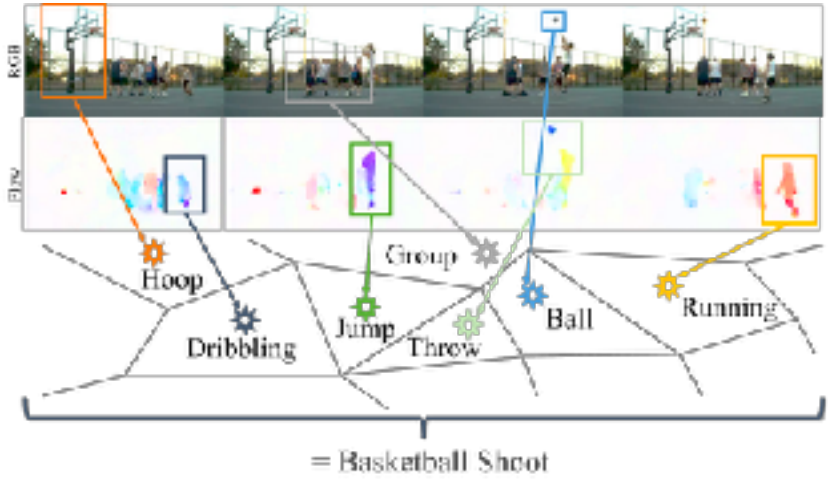

Teaching Assistants

Senthil Purushwalkam

First year PhD in Robotics

Research Interests:

Unsupervised, weakly supervised and web supervised learning

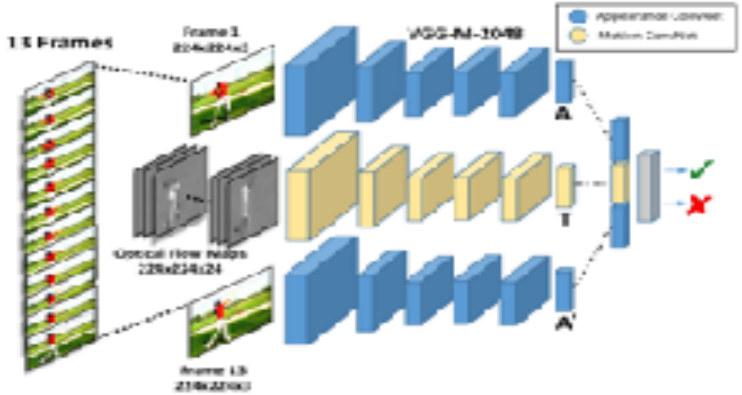

Pose from Action: Unsupervised Learning of Pose Features based on Motion, AAVL Workshop - ECCV 16

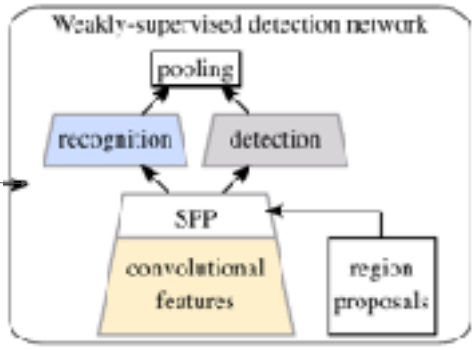

Teaching Assistants

Nai Chen Chang

Second year MS in Robotics

Research Interests:

Representations in Human Vision vs. Deep Learning

16-824: Learning-based Methods in Vision

What is this course about?

What is the goal of Computer Vision?

Systems that can "understand" Visual Data

understanding visual data

understanding visual data

understanding visual data

What does it mean to understand?

The Vision Story Begins…

“What does it mean, to see? The plain man's answer (and Aristotle's, too). would be, to know what is where by looking.”

-- David Marr, *Vision (1982)*

Slide Credit: Alyosha Efros

Vision: a split personality

"What does it mean, to see? The plain man's answer (and Aristotle's, too). would be, to know what is where by looking. In other words, vision is the process of discovering from images what is present in the world, and where it is."

Answer #2: looks like flat sittable surface of the couch

Which do we want?
Is the difference just a matter of scale or is there some fundamental difference?

Answer #1: *pixel of brightness 243 at position (124,54) …and depth .7 meters*

two-sides of vision

Measurement vs. Perception

Brightness: Measurement vs. Perception

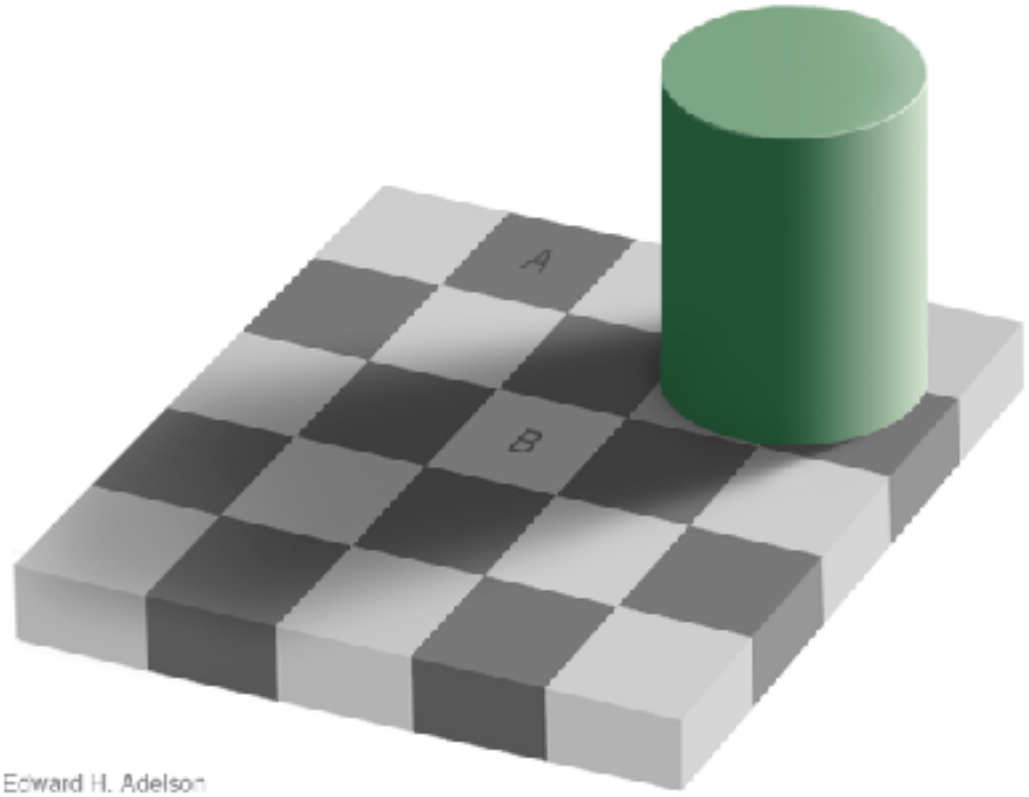

Slide Credit: Alyosha

Brightness: Measurement vs. Perception

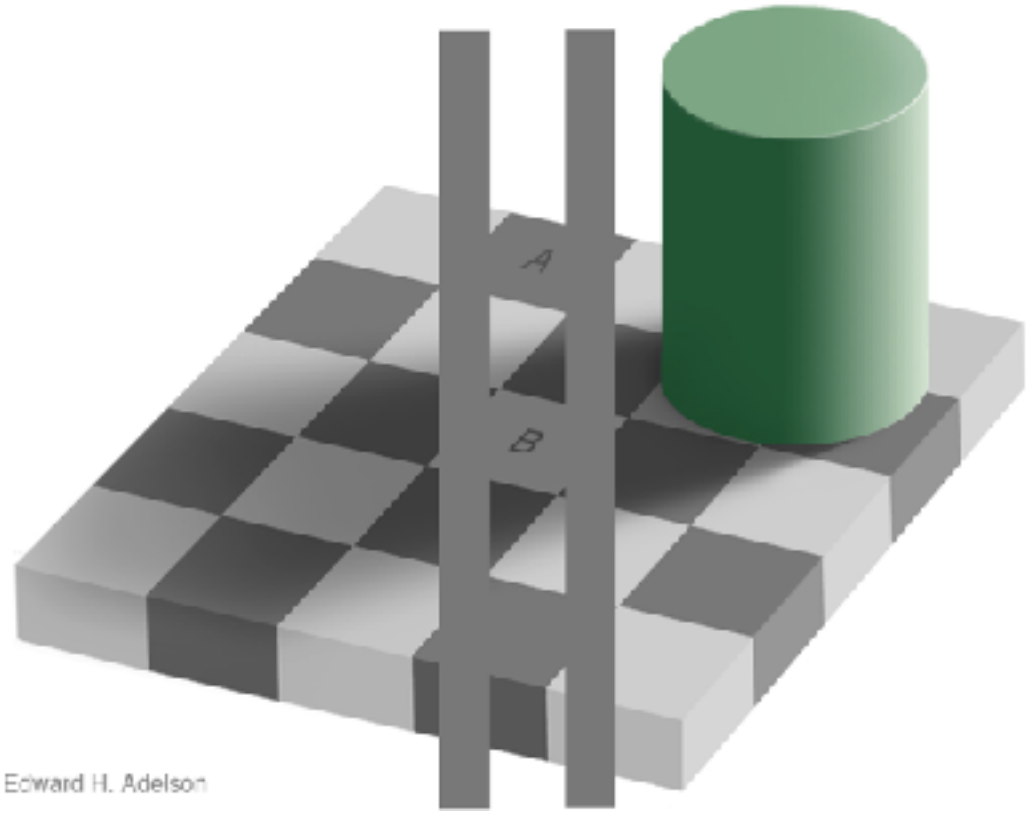

Proof!

Slide Credit: Alyosha

Measurement

Length

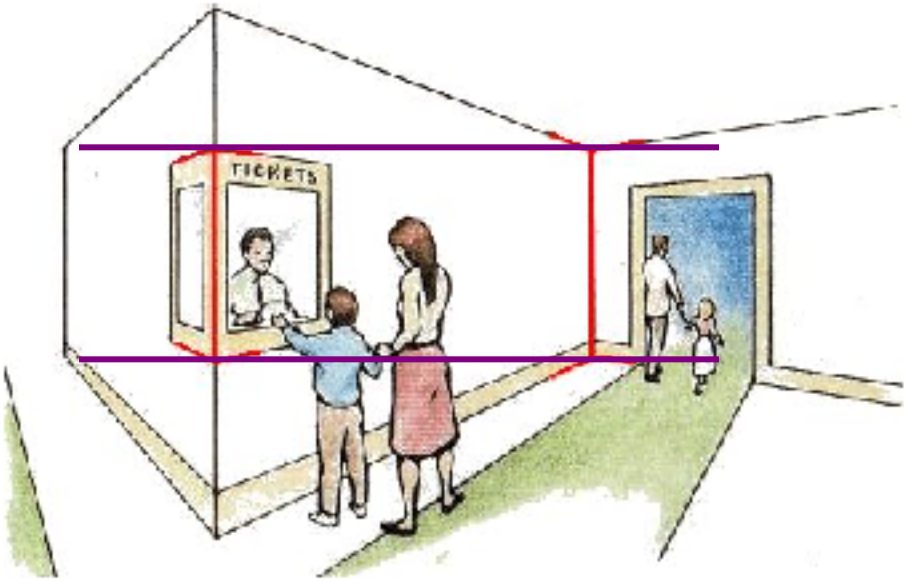

Müller-Lyer Illusion

http://www.michaelbach.de/ot/sze_muelue/index.html

Slide Credit: Alyosha

Measurement

Capturing physical quantities like pixel brightness, depth, etc.

Perception/Understanding

- a high-level representation that captures the semantic structure of the scene and its constituent objects.
- Subjective – Depends on Task and Agent
- Intersection of what you see and what you believe (prior knowledge)

Vision as Measurement Device

Real-time stereo on Mars

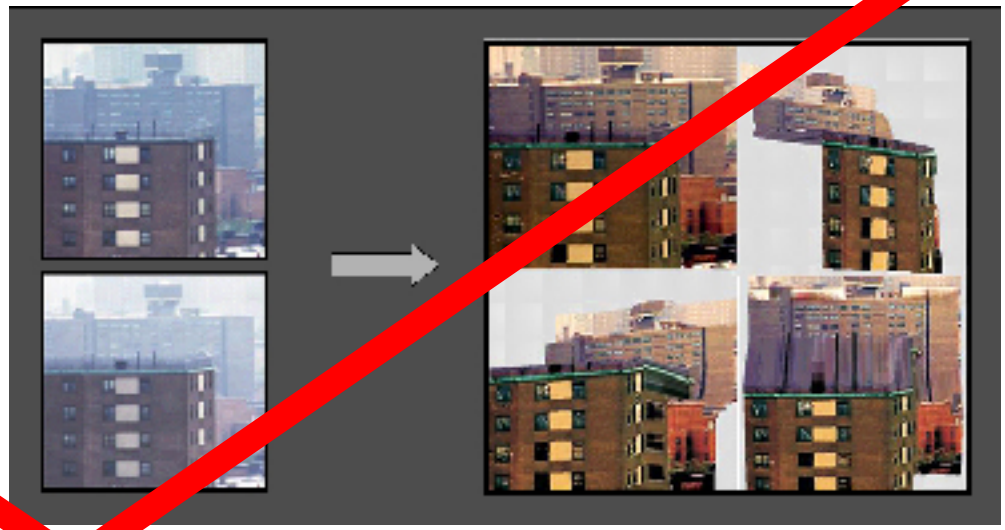

Physics-based Vision

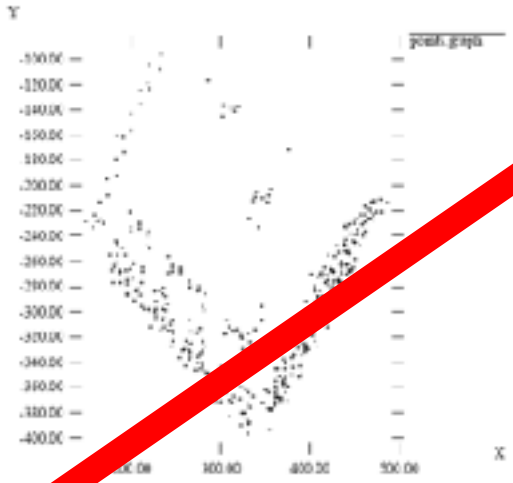

Structure from Motion

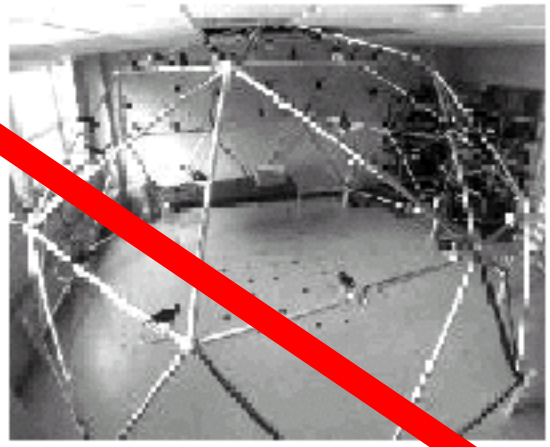

Virtualized Reality

Slide Credit: Alyosha Efros

…but why do we care about perception?

The goals of computer vision (**what + where**) are in terms of what humans care about.

So what do humans care about?

Image Classification/ Scene Recognition

Object Detection

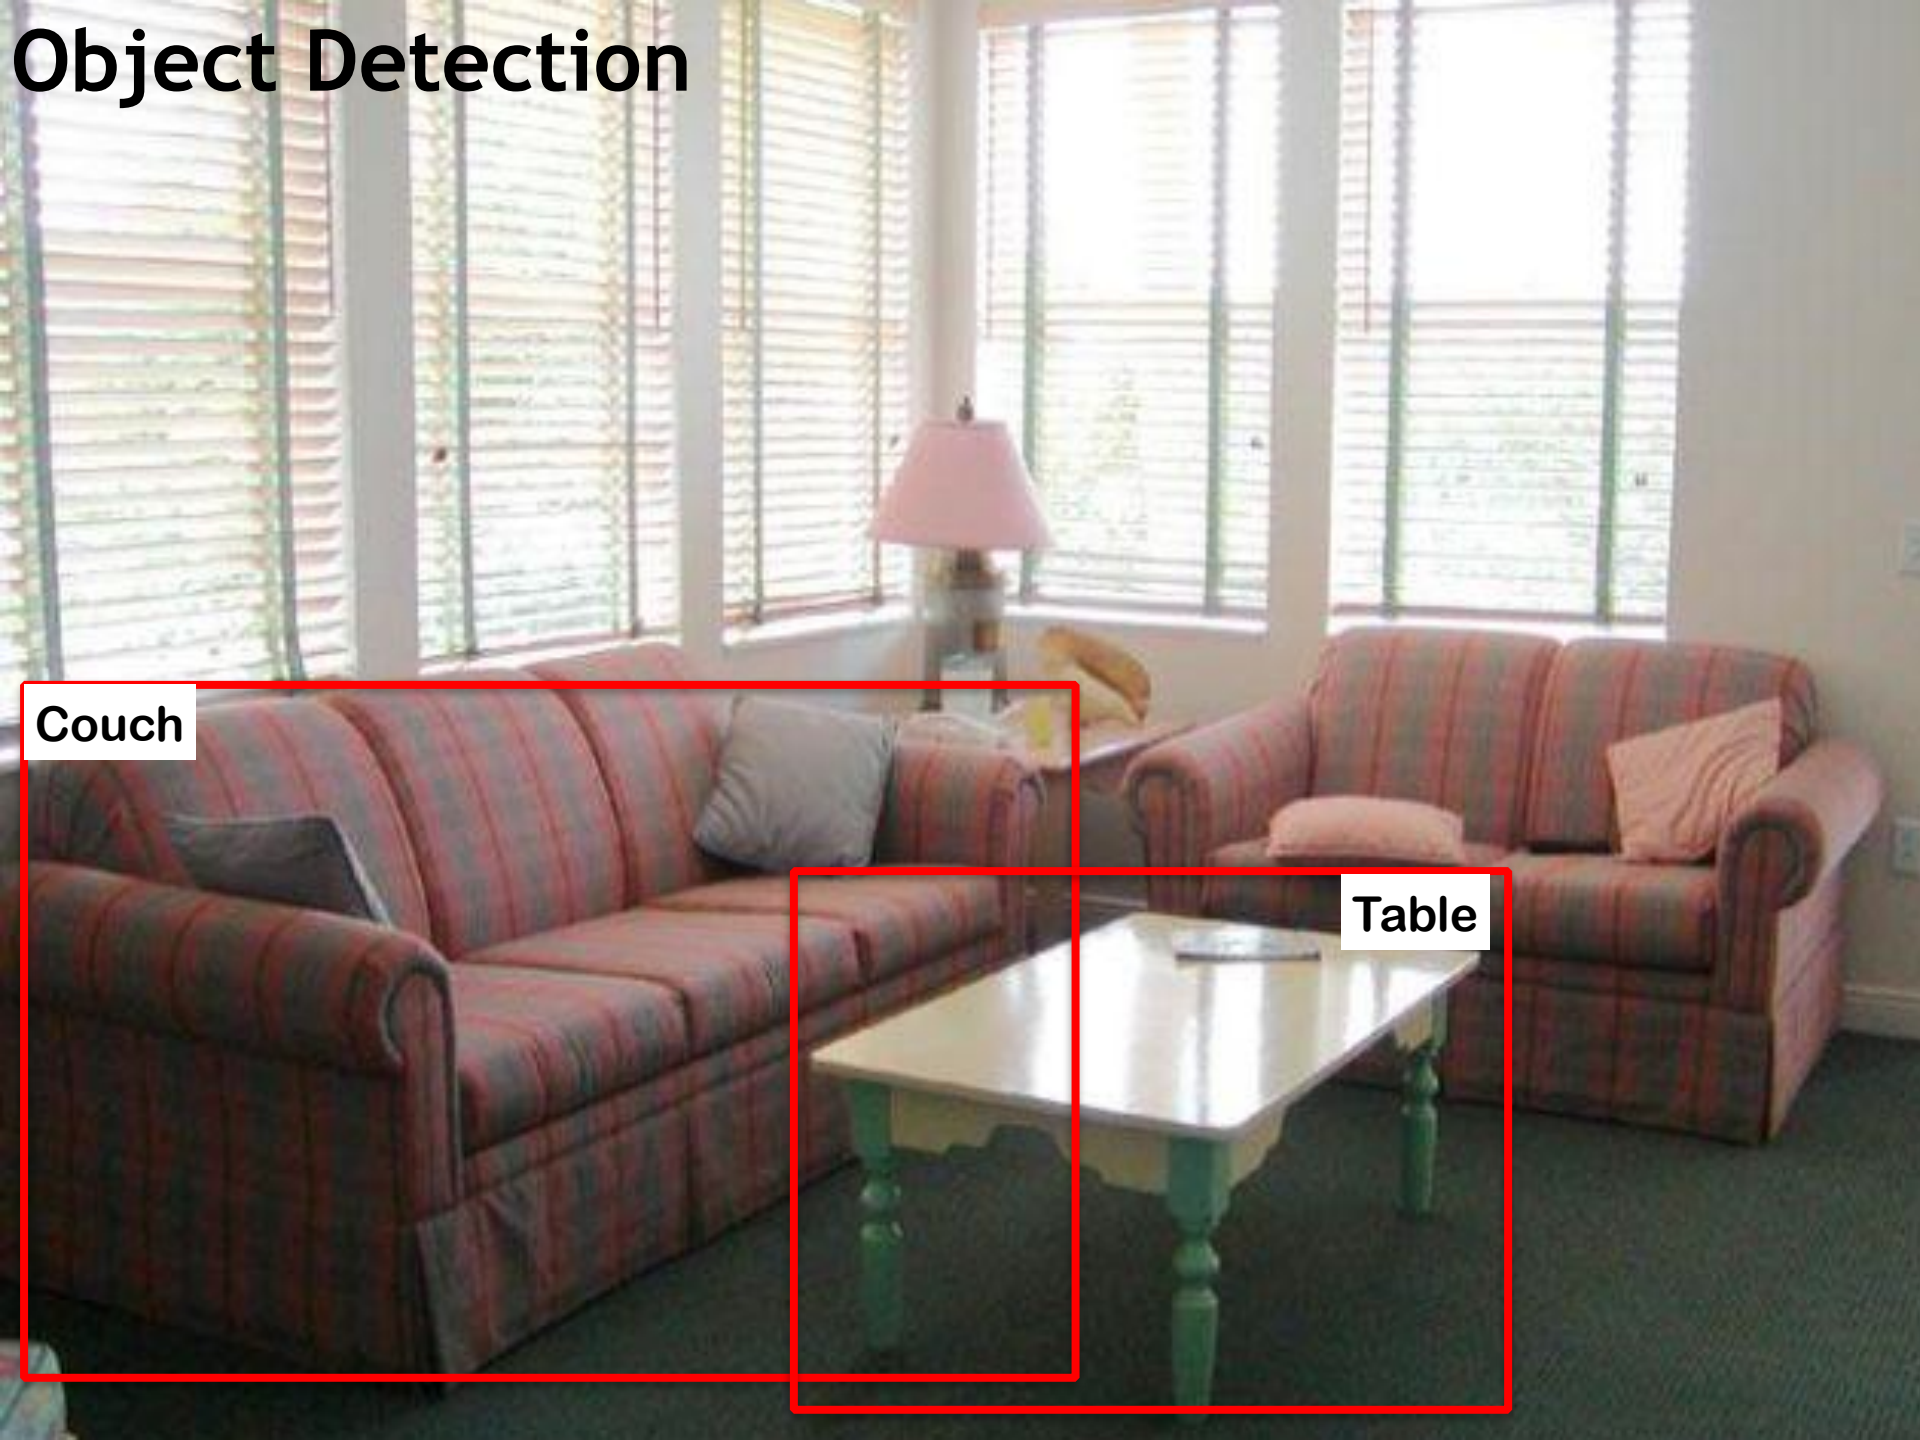

Object Segmentation/Categorization

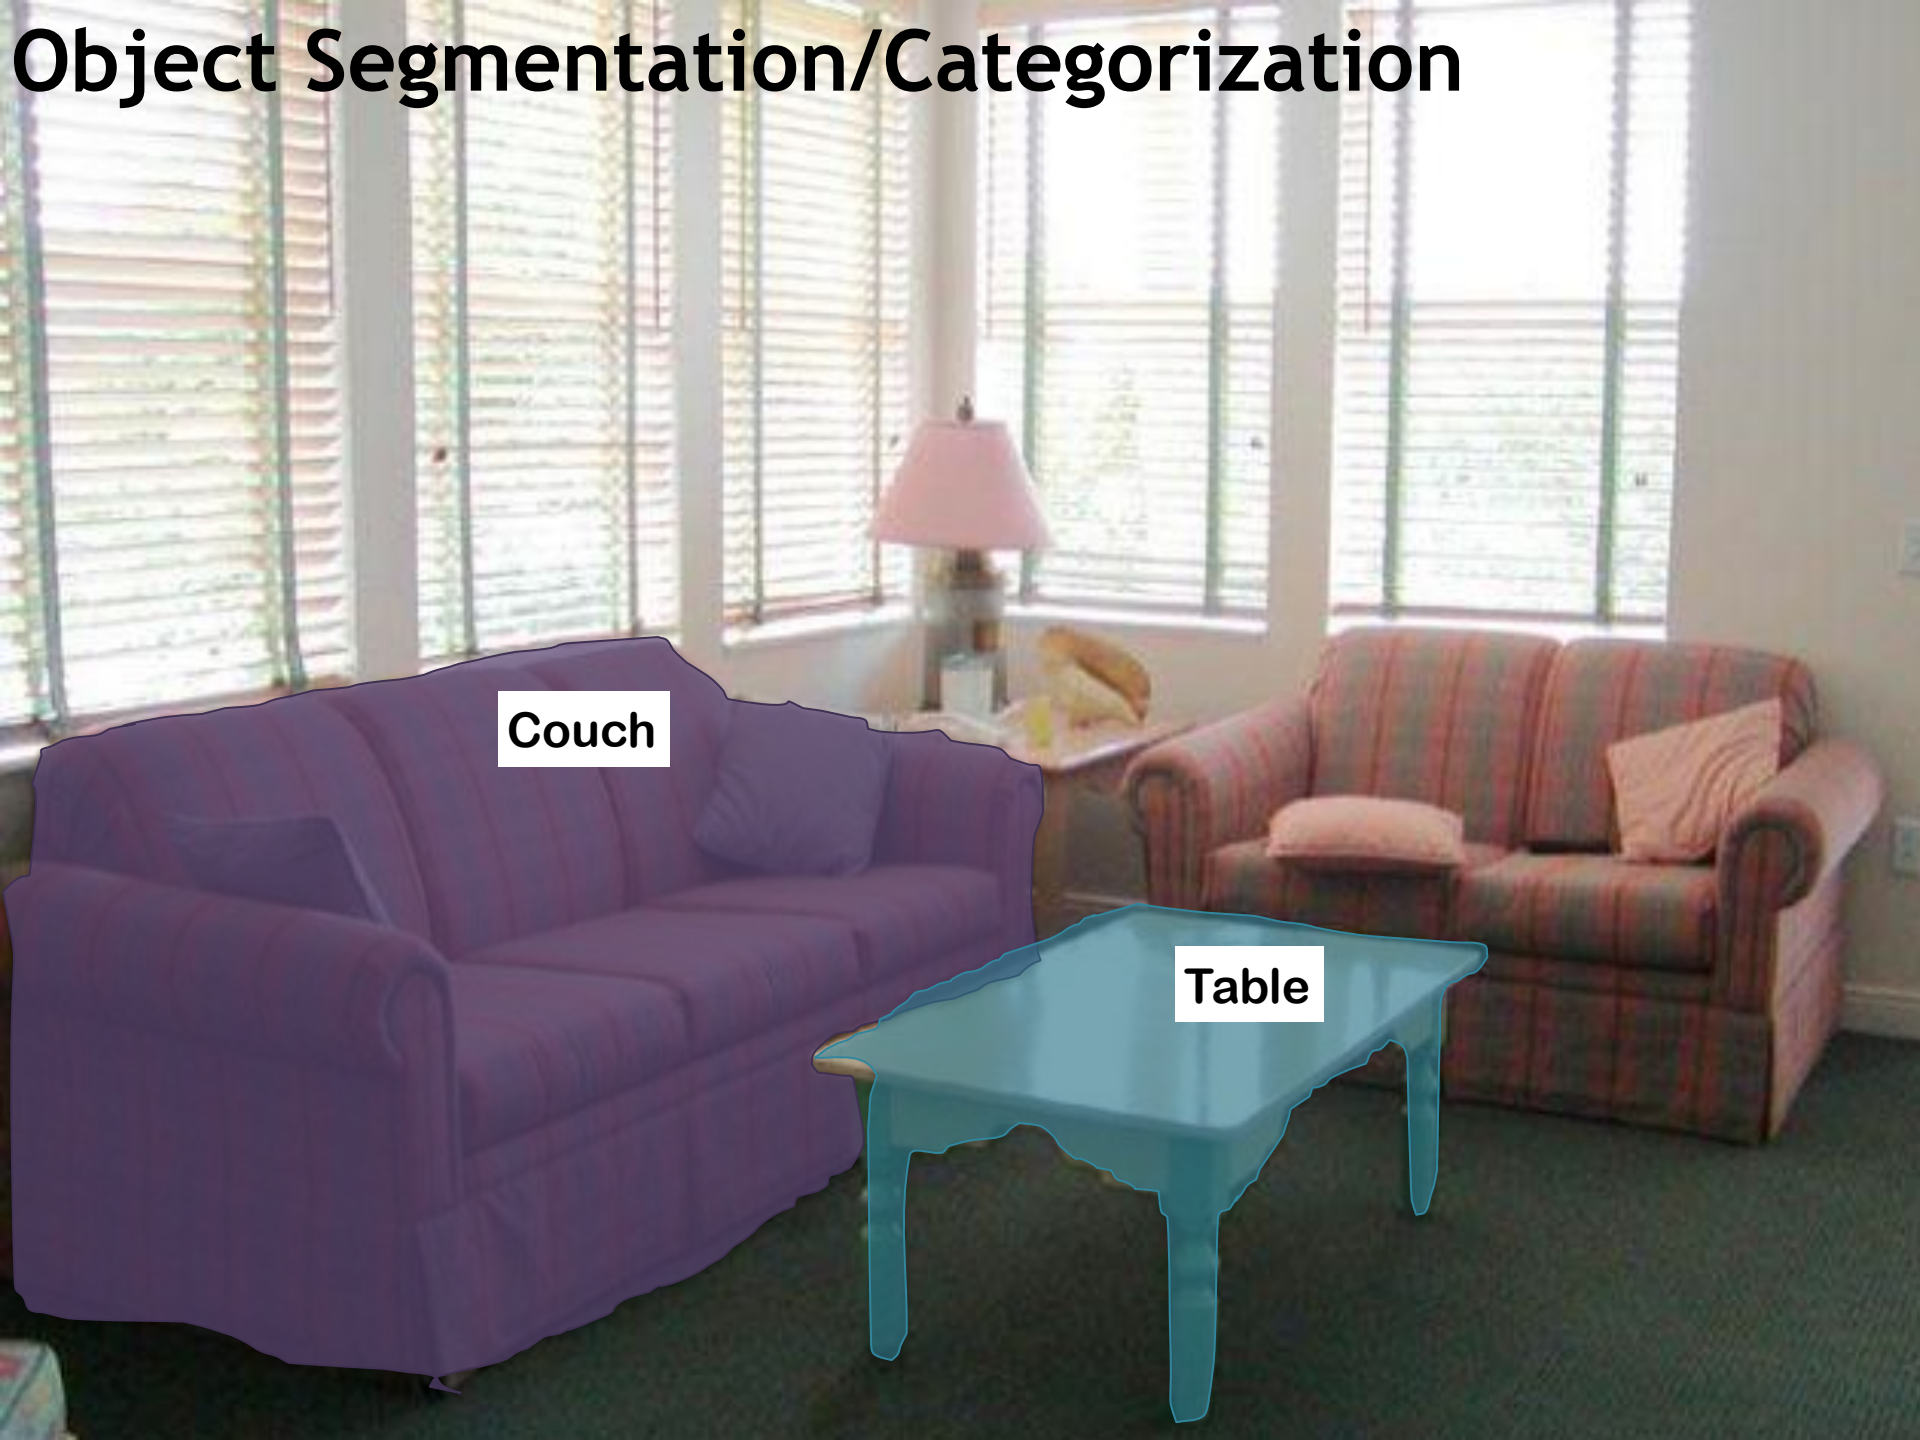

3D Understanding

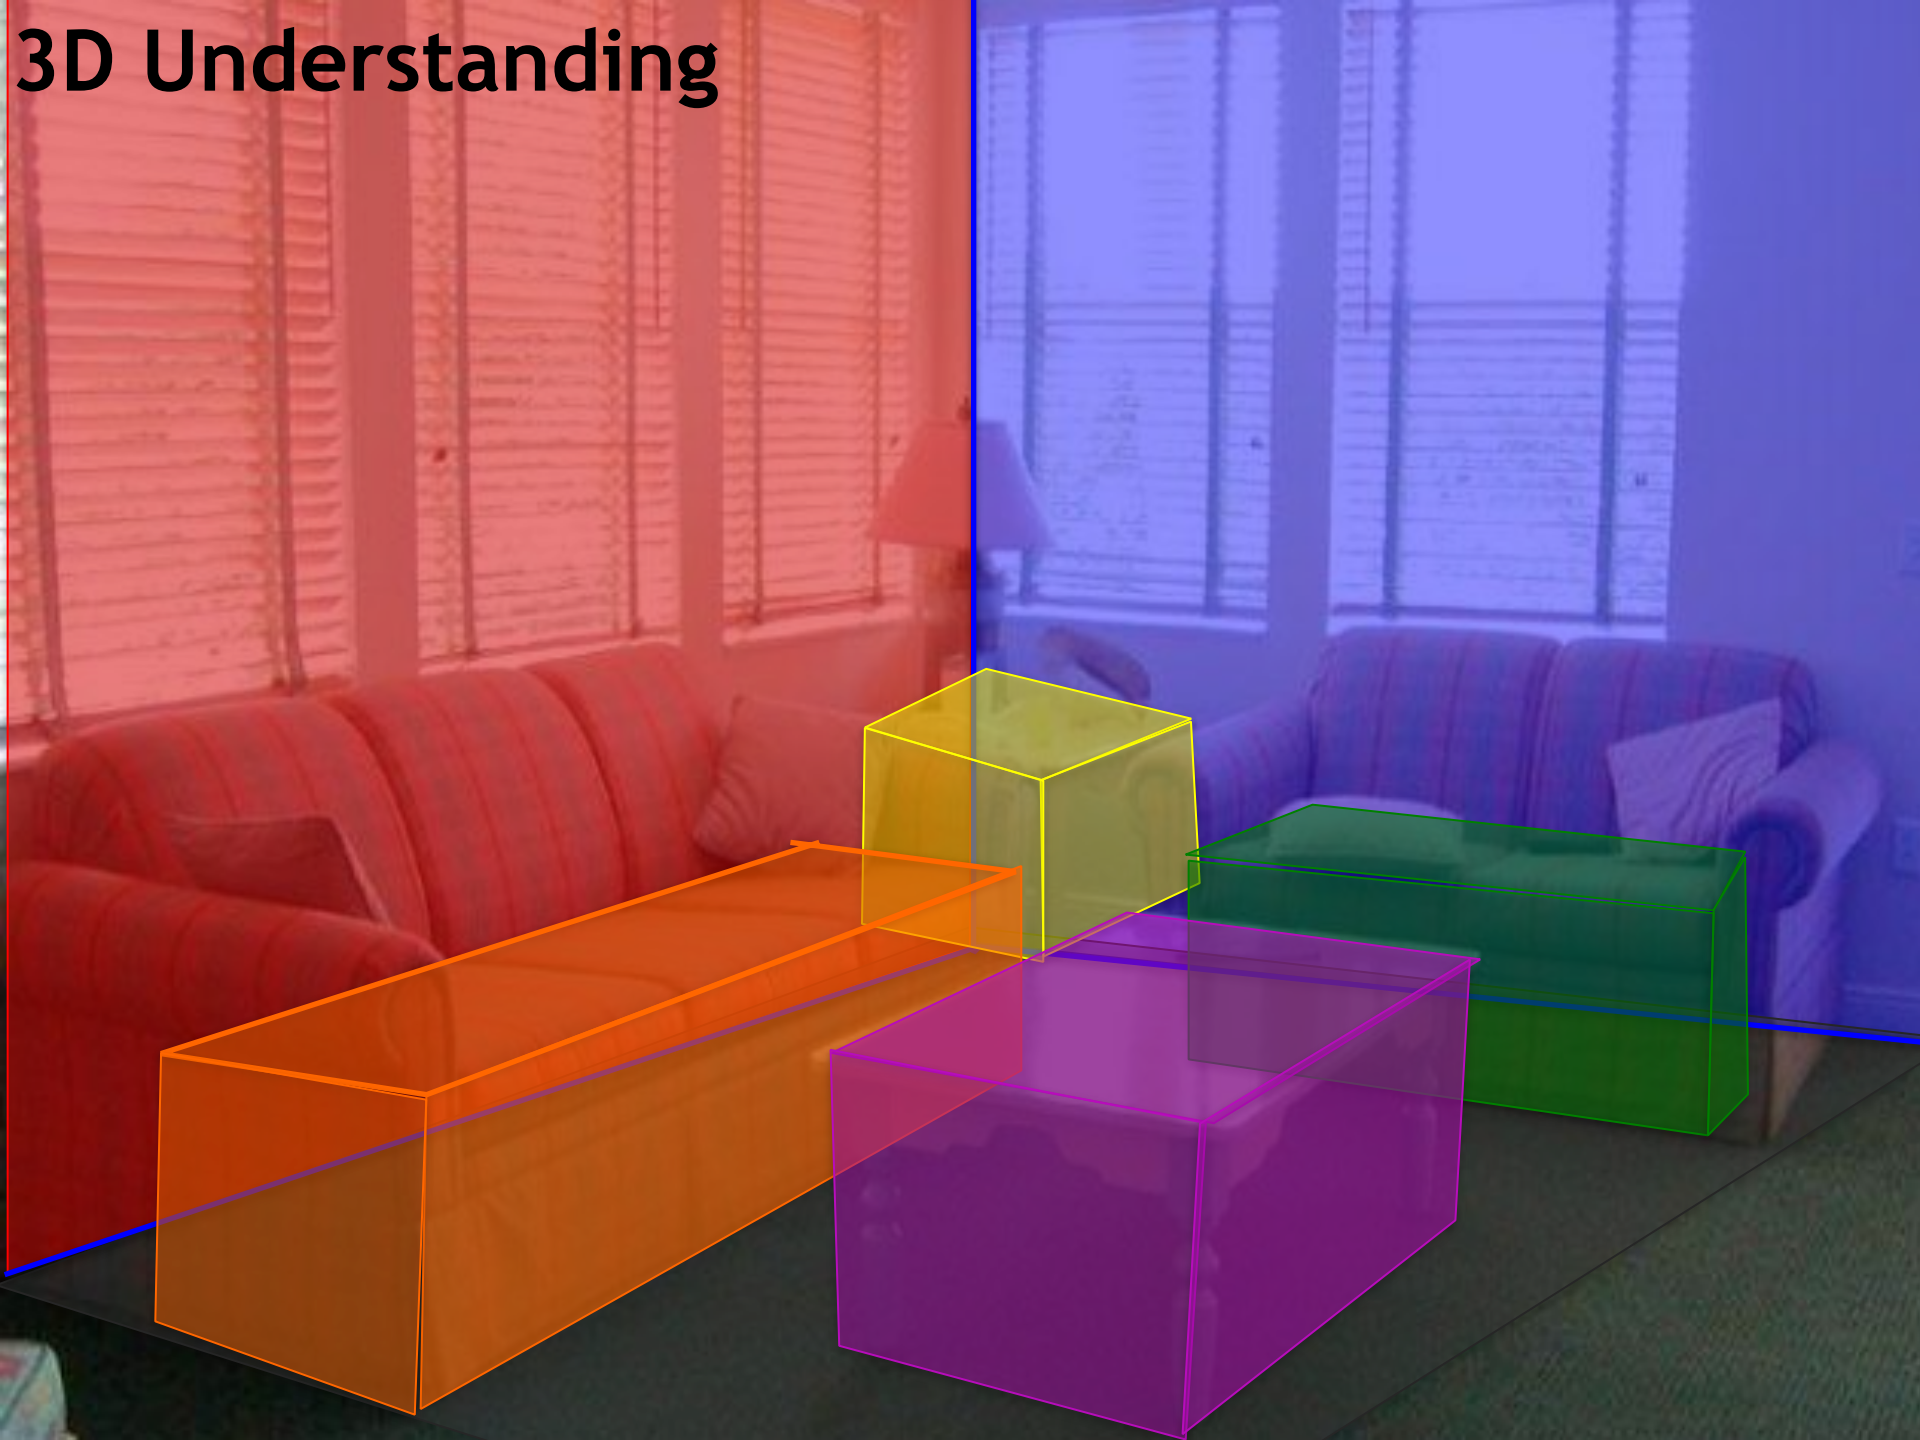

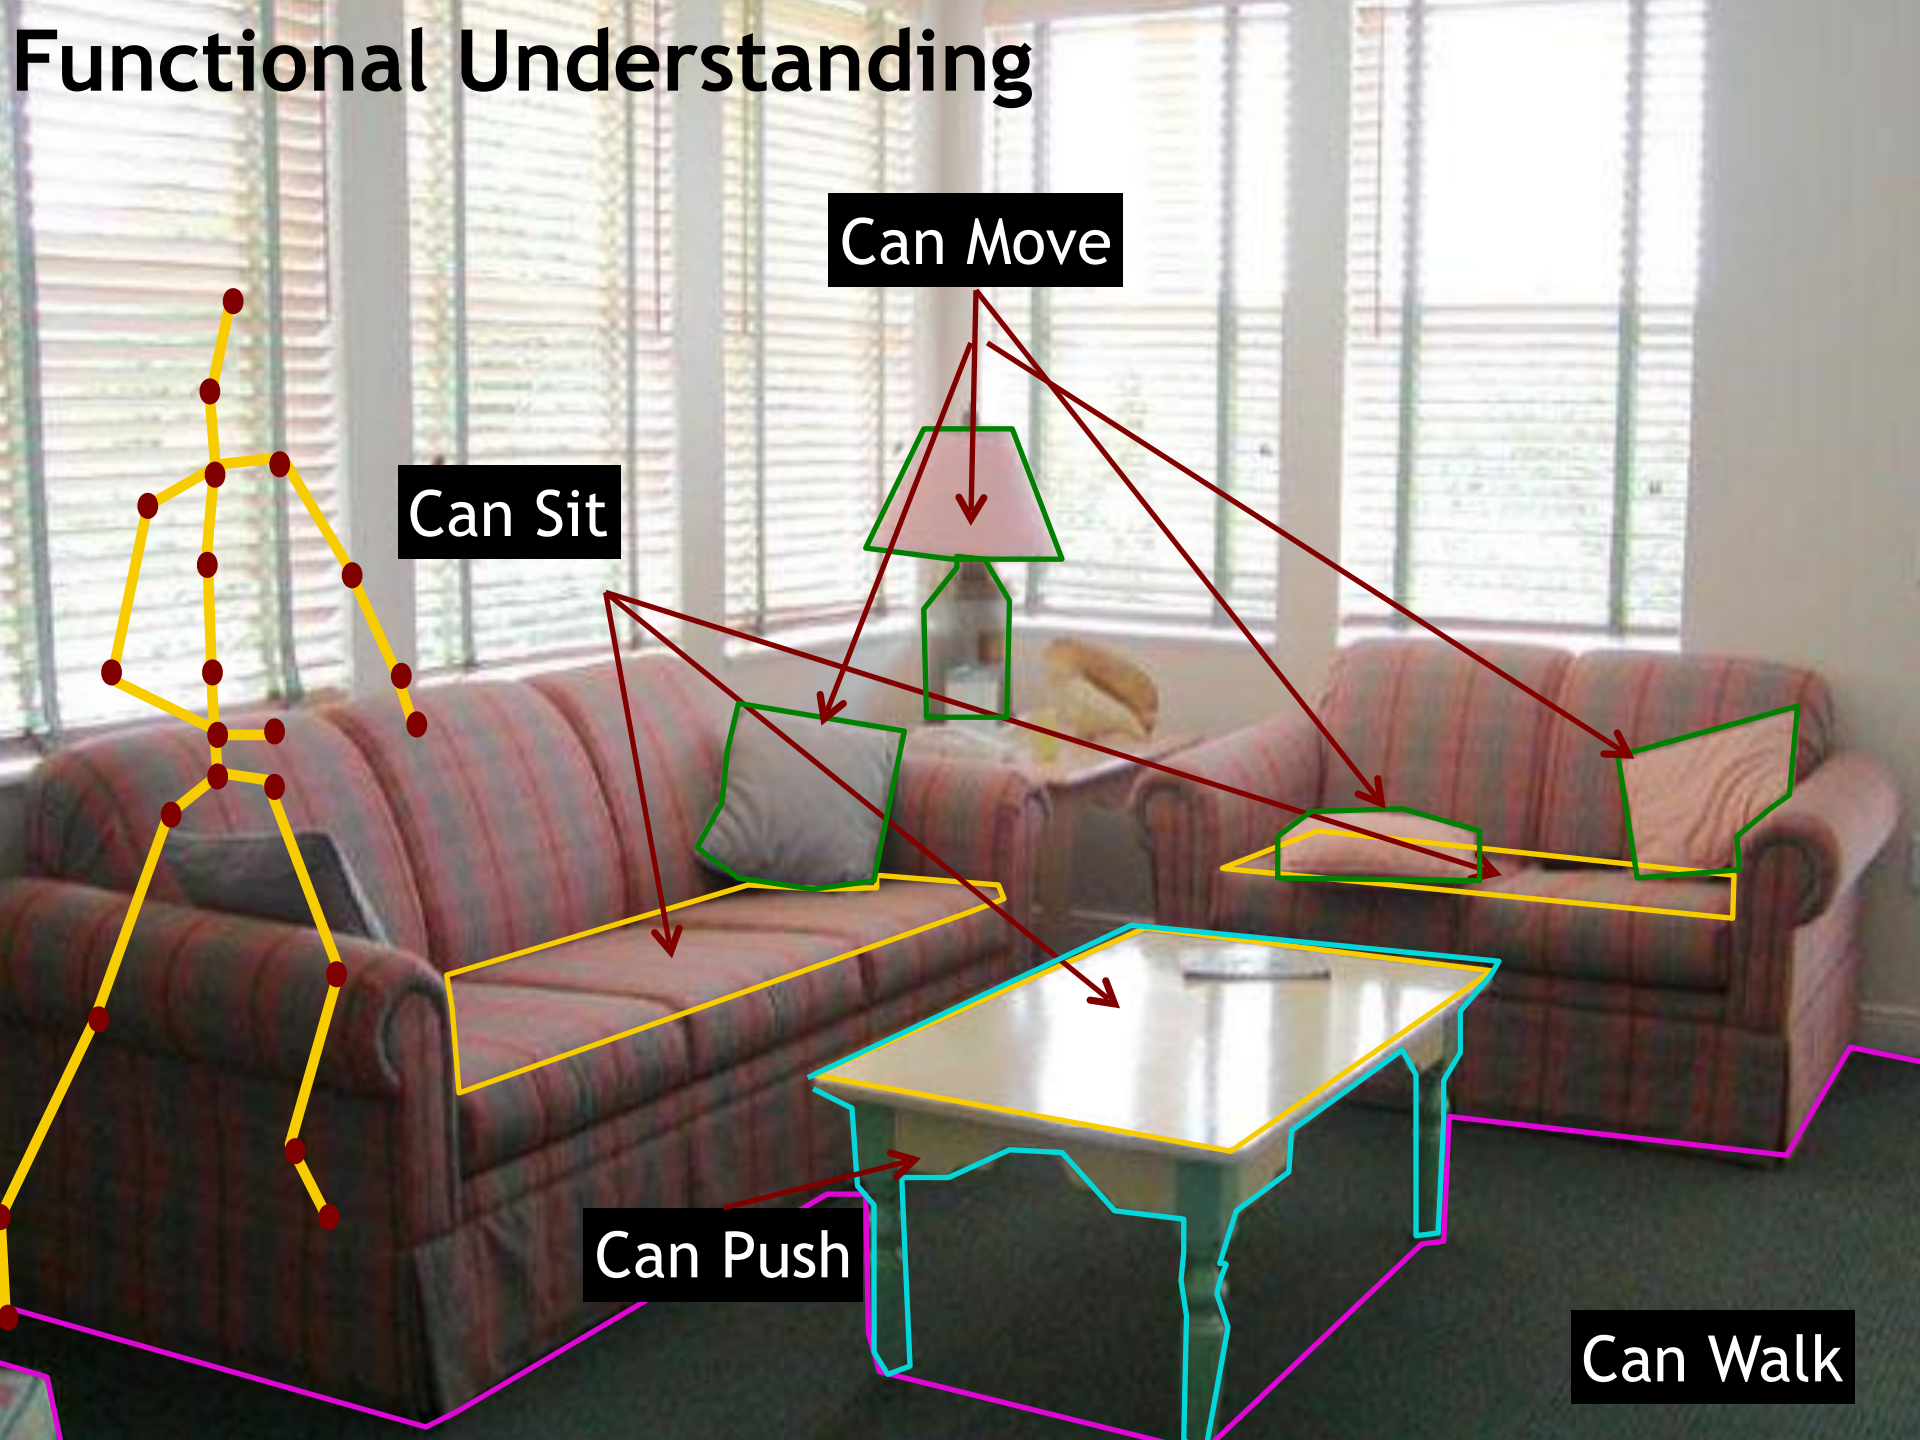
Functional Understanding
Can Move
Can Sit
Can Push
Can Walk

Pose Estimation:

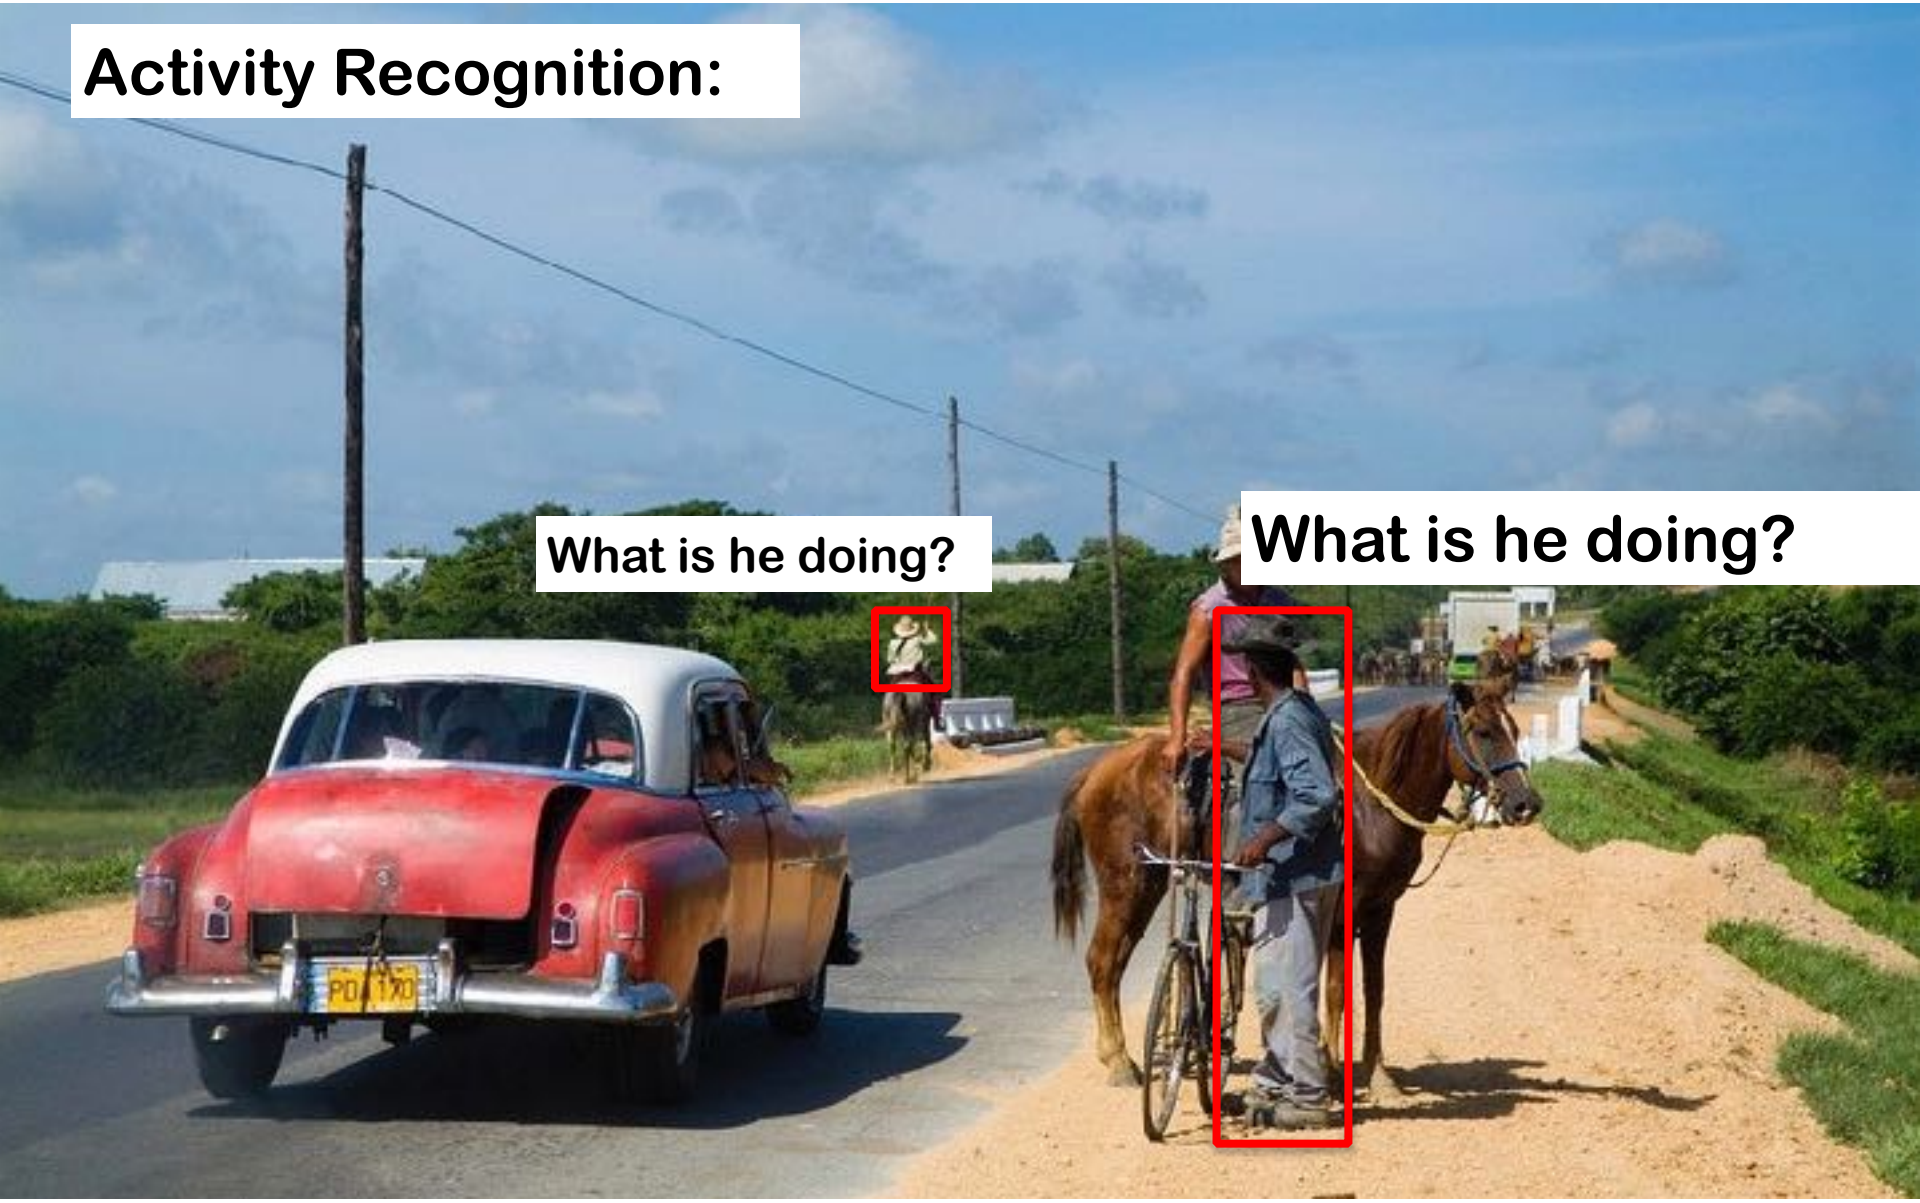
Activity Recognition:
What is he doing?
What is he doing?

Why are these problems hard?

Challenges 1: view point variation

Michelangelo 1475-1564

slide by Fei Fei, Fergus & Torralba

Challenges 2: illumination

slide credit: S. Ullman

Challenges 3: occlusion

Magritte, 1957

slide by Fei Fei, Fergus & Torralba

Challenges 4: scale

slide by Fei Fei, Fergus & Torralba

Challenges 5: deformation

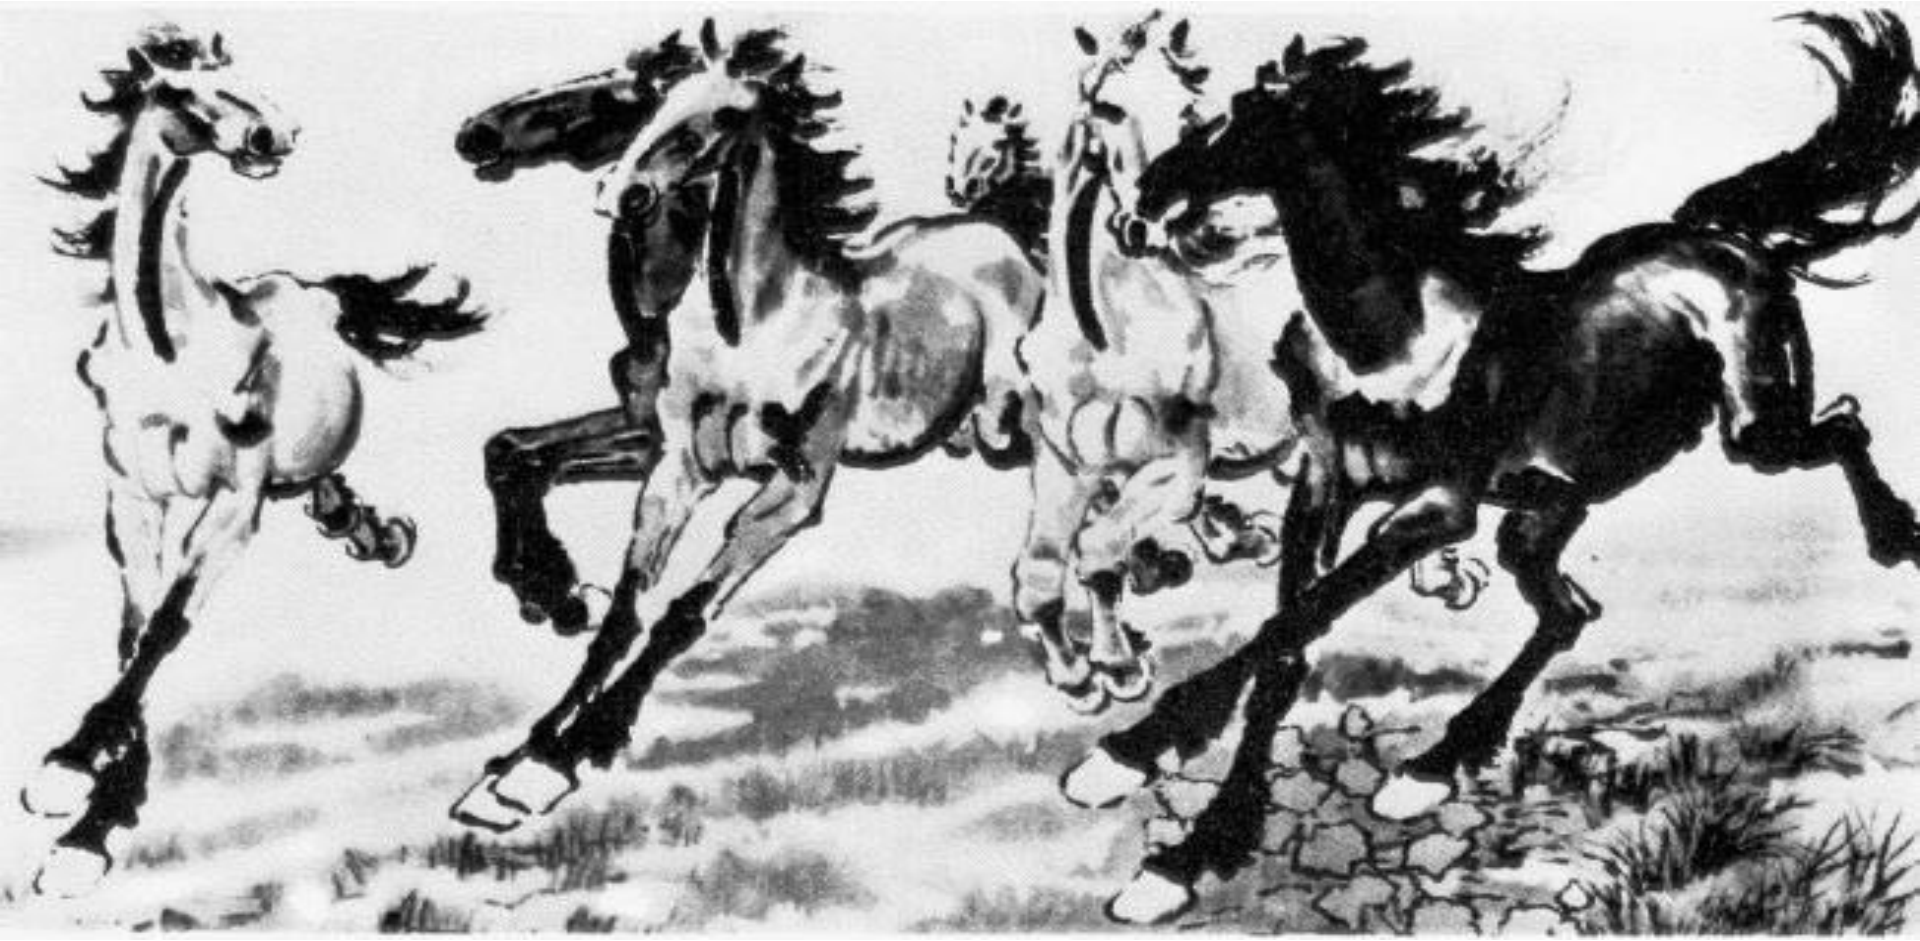

Xu, Beihong 1943

Challenges 6: background clutter

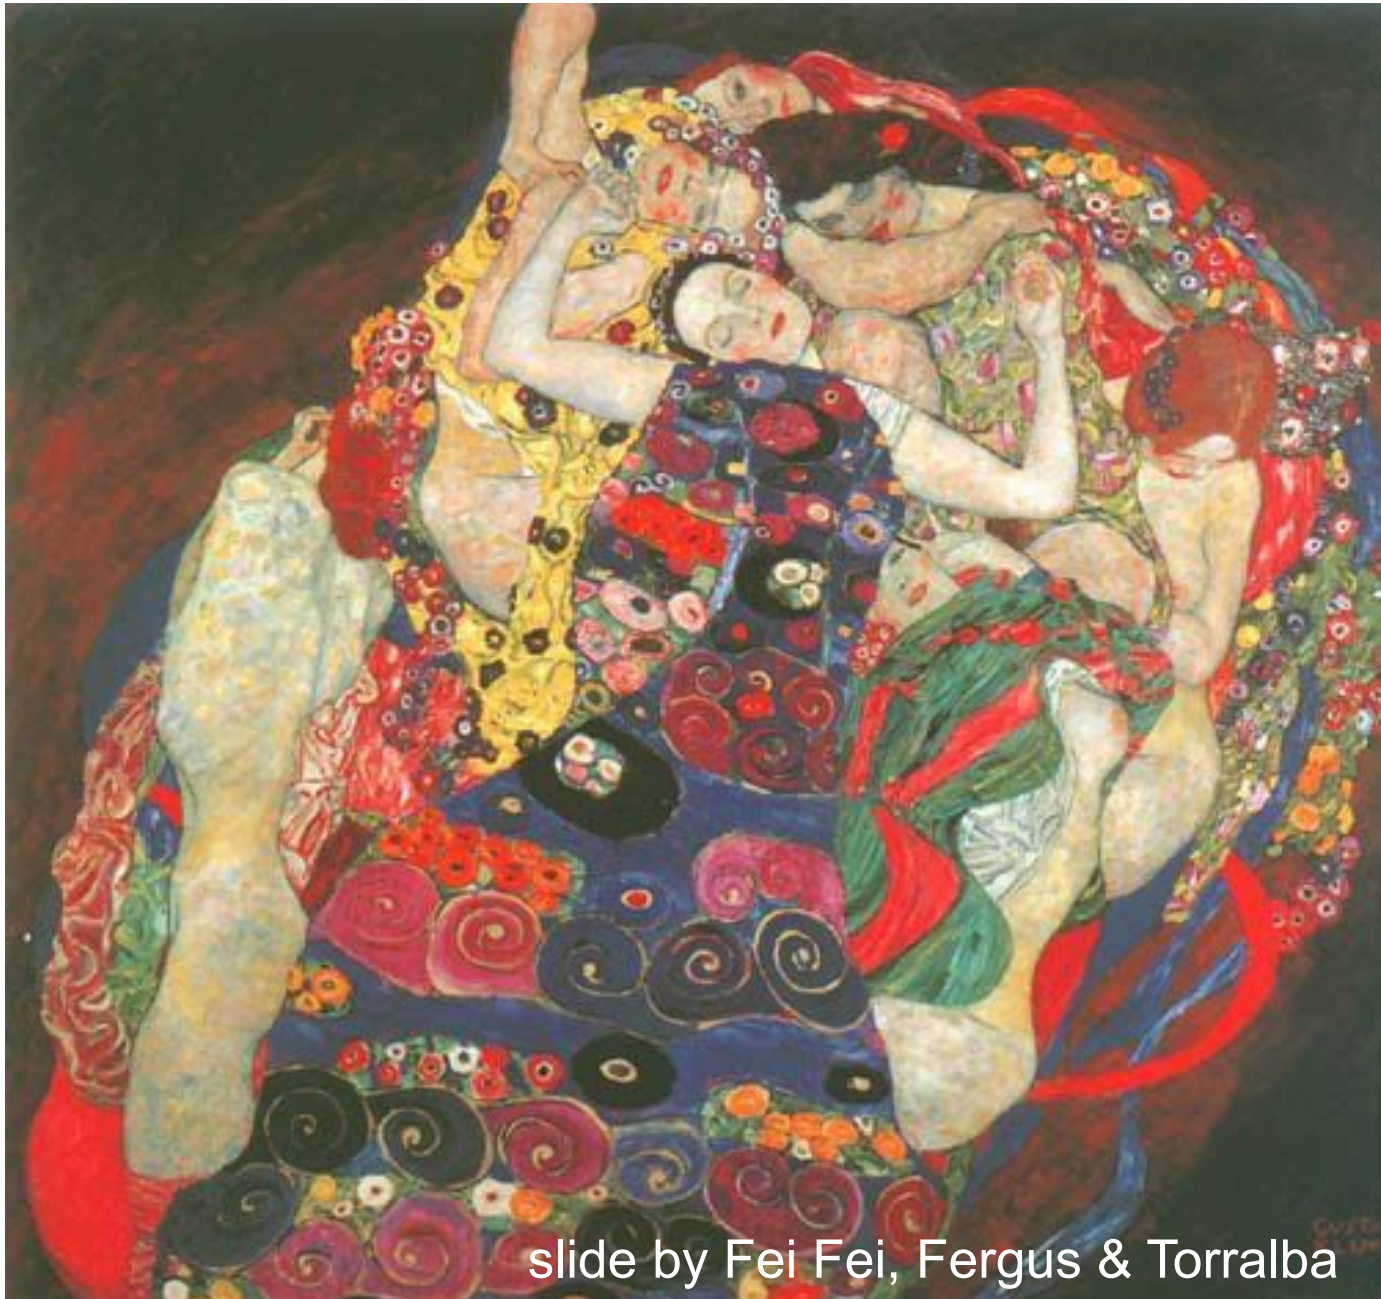

Klimt, 1913

slide by Fei Fei, Fergus & Torralba

Challenges 7: object intra-class variation






slide by Fei-Fei, Fergus & Torralba

Challenges 8: local ambiguity

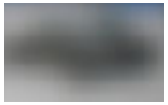

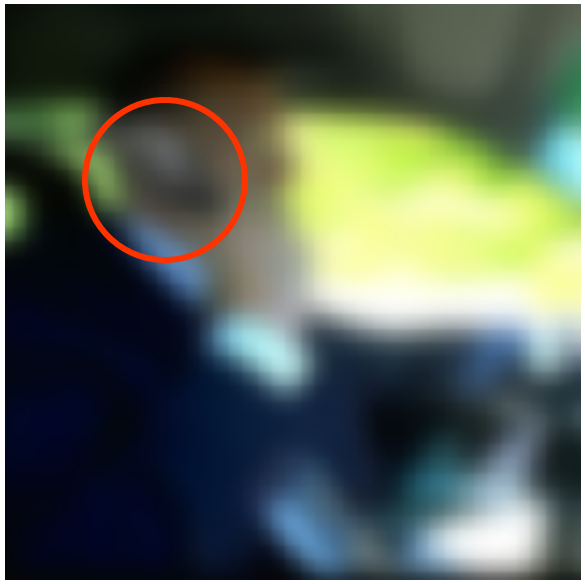

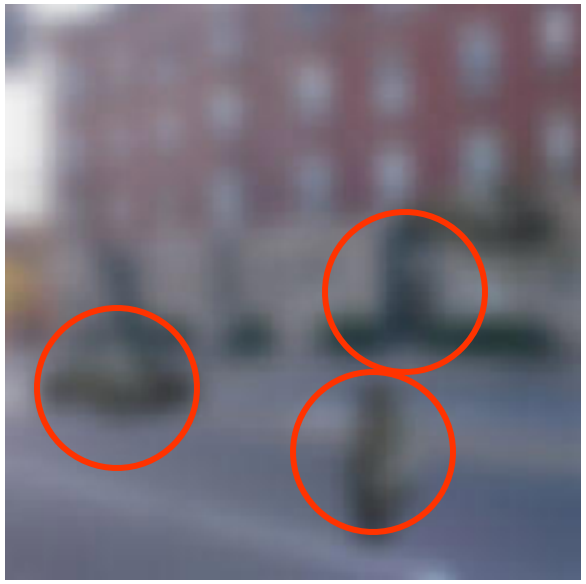

slide by Fei-Fei, Fergus & Torralba

Challenges 9: the world behind the image

ill-posed

- EXAMPLE:
- Recovering 3D geometry from **single** 2D projection

- Infinite number of possible solutions!

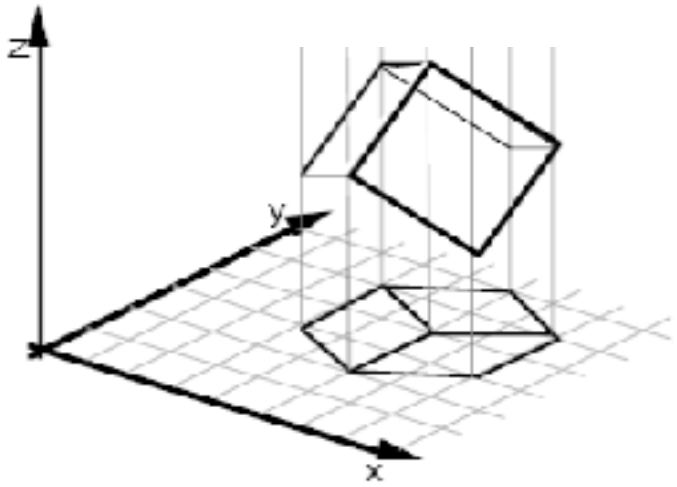

from [Sinha and Adelson 1993]

How do we solve it?

Data to Rescue !!

- Data to build observation models..
- Data to build priors about the visual world.
- Use the models and prior information to infer..

Machine-Learning!

In this course, we will:

Take a few baby steps…

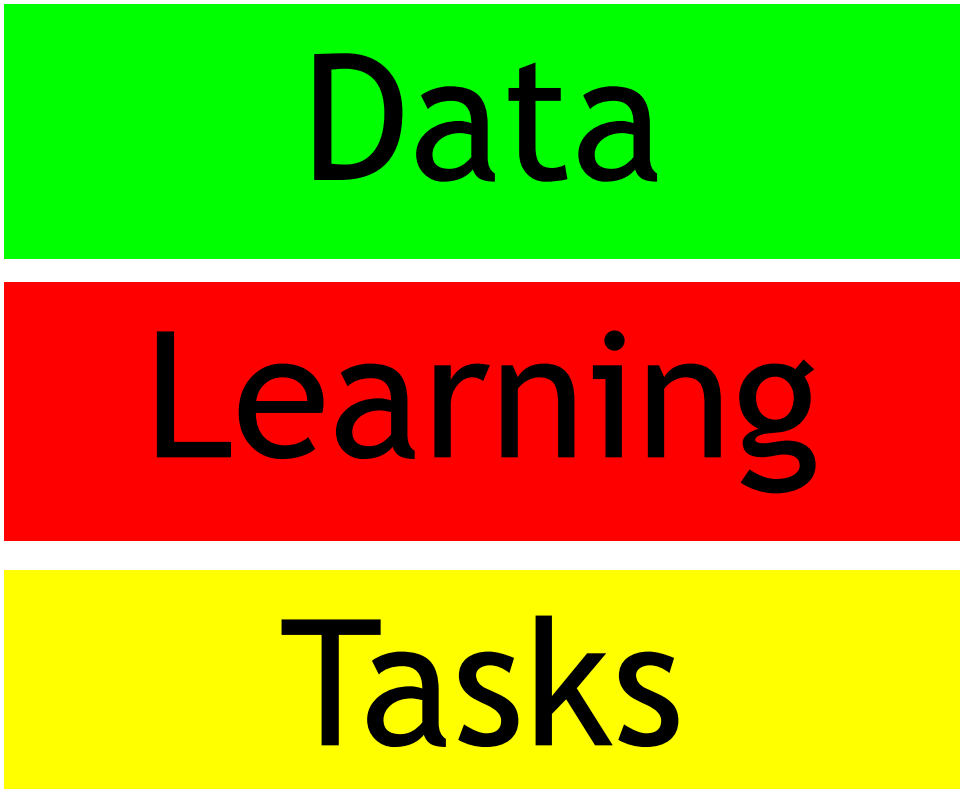
Data
Learning
Tasks

Technical Challenges

Challenges: scale, efficiency

6 billion images

70 billion images

1 billion images served daily

100 hours uploaded per minute

10 billion images

From 

Almost 90% of web traffic is visual!

Challenges: scale, efficiency

Fei Fei Li, Rob Fergus, Antonio Torralba

Biederman 1987

Challenges: learning with minimal supervision

Kristen Grauman

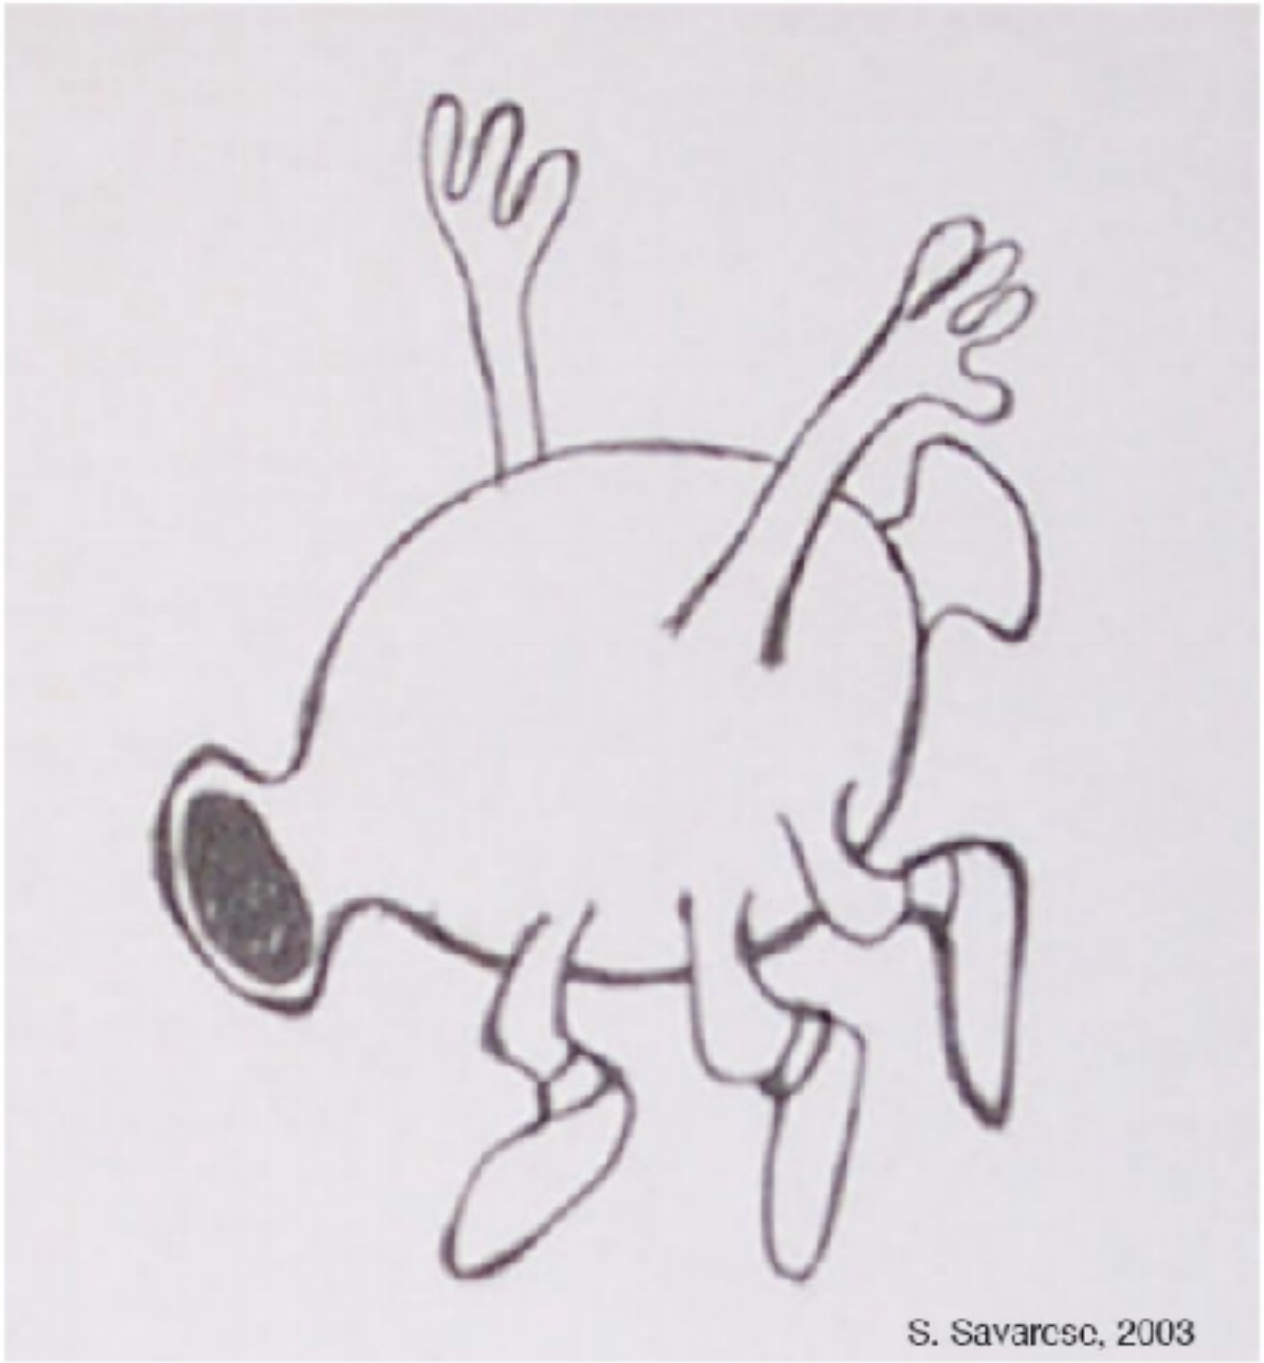

This is a pottopod

Slice from Pietro Perona, 2004 Object Recognition workshop

Find the pottopod

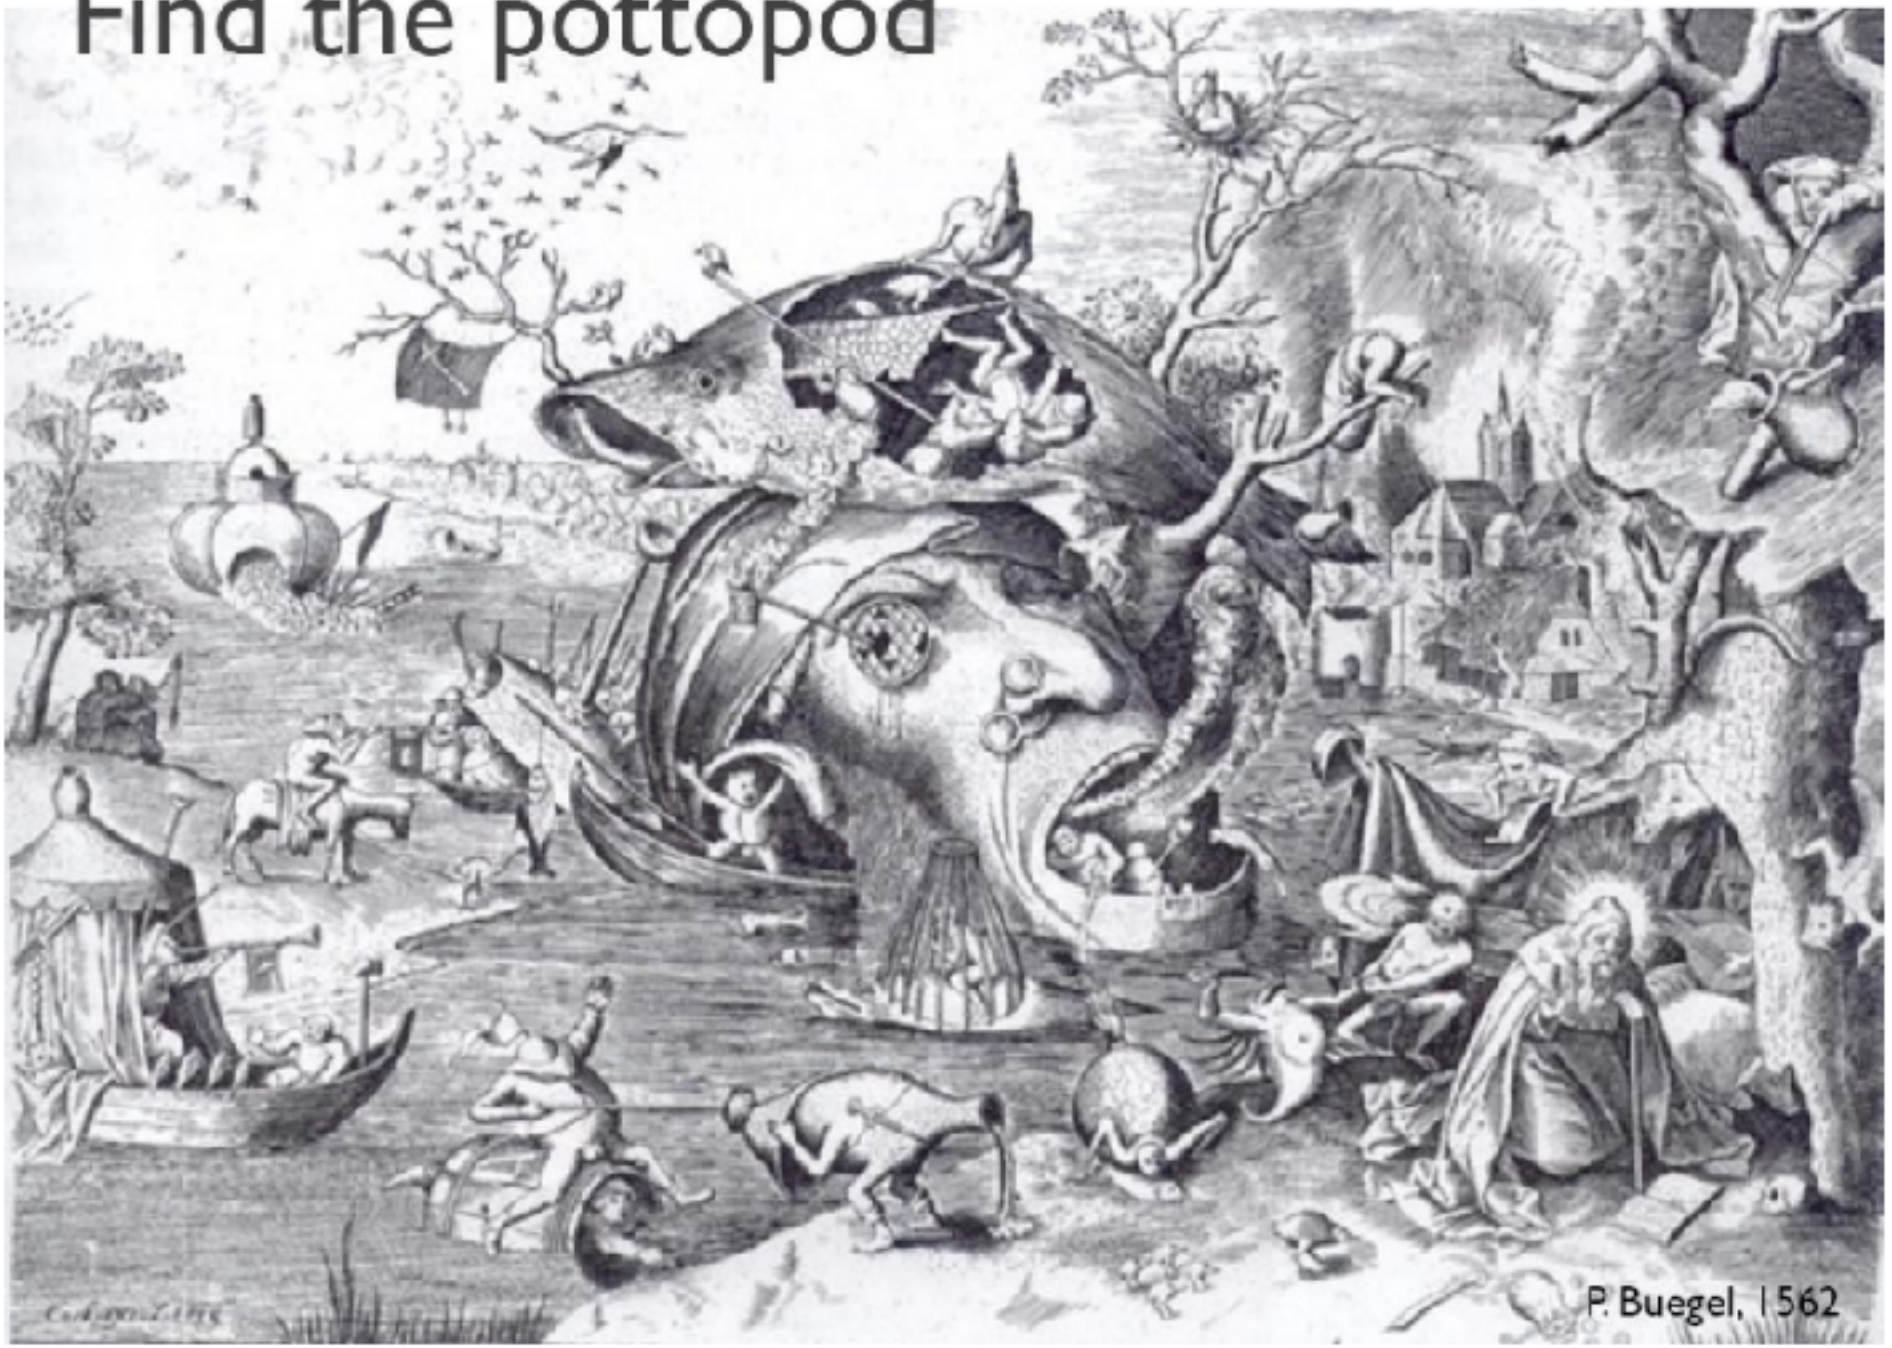

Slide from Pietro Perona, 2004 Object Recognition workshop

What to expect in the class?

Learning as a tool to exploit big data, build prior models etc.

Not formulate problem in complicated manner…

Learning as a tool to exploit big data, build prior models etc.

Not formulate problem in complicated manner…

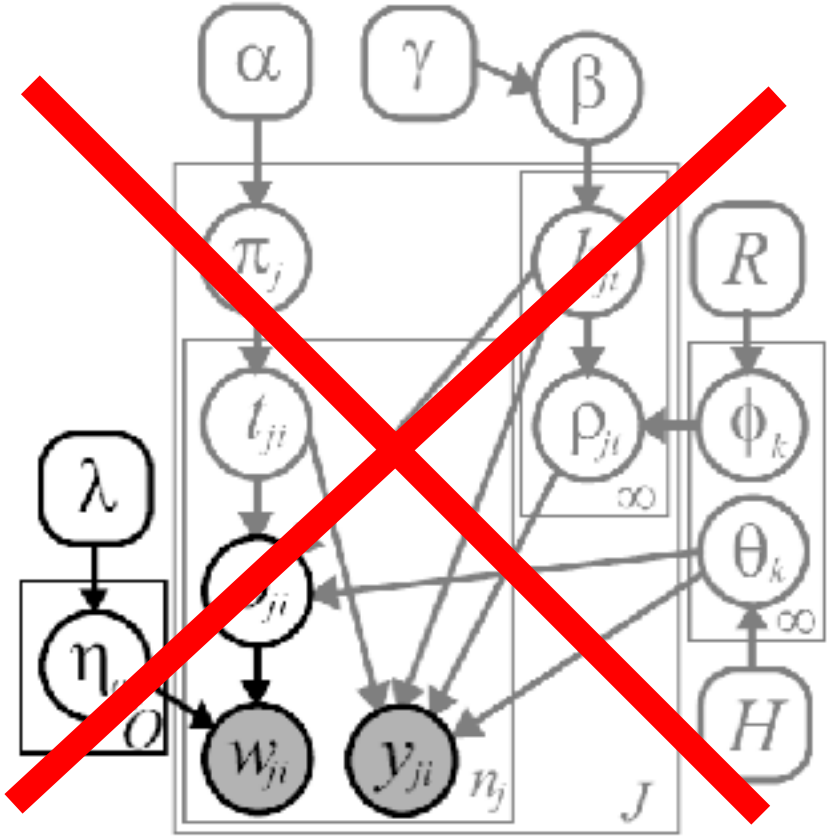

But that said…

- We will still look at the learning methods which give the state of the art performance on these tasks.

- For example, most focus this year will be on deep learning – Convolutional Neural Networks (CNN)..

Is this a research course?

- Two year ago – YES!
- But times have changed: Computer Vision is a hot topic in industry now..
- 2012 – Resurgence of Deep Networks (CNNs)

2014 – Deep Learning is Everywhere

- Google, Facebook, Baidu, Apple
  - Strong deep learning groups hiring everywhere..
  - Beyond Research: Development
    - Image Search
    - Automated Driving

Startups Sold Everyday

- Vision Factory, EuVision, Flutter….

Course Outline

Goals

- Develop background for Convolutional Neural Networks.
- Know the established state of the art in each task areas.
- Get up to speed on big chunk of vision research
  - understand 70% of CVPR papers!
- Read some interesting papers together
  - Learn something new: both you and us!
- Develop background and expertise to use ConvNets in your own work
- Learn how to speak
- Learn how think critically about papers

Course Organization

- Requirements:
  1. Class Participation (15%)
     - Keep annotated bibliography
     - Post on the Class Blog before **each class**
     - Ask questions / debate / flight / be involved!
  2. Project (25%)
  3. Assignment (3x20%)

Class Participation

- Keep annotated bibliography of papers you read (always a good idea!). The format is up to you. At least, it needs to have:
  - Summary of key points
  - A few Interesting insights, “aha moments”, keen observations, etc.
  - Weaknesses of approach. Unanswered questions. Areas of further investigation, improvement.

- Submit a comment on the Class Blog
  - ask a question, answer a question, post your thoughts, praise, criticism, start a discussion, etc.

Class Assignment

Three assignments to get you familiar with deep learning.

Toolboxes

- PYTORCH
- Tensor flow

Class Project

Opportunity to work on the crazy idea which your advisor would not let you do ! (Group of 2-3)

Merit Criteria

1. Crazy (the more different it sounds the better it is)
2. Amount of Work/Results.
3. Report/Presentation

Failure/Success has no points! An idea with interesting failure results is a successful project!

Logistics

- Huge Waitlist
- Talk to me after class!