


The Promise and Perils of Big Data




Some Slides from A. Efros and A. Torralba

Why do we need data?

Most problems in vision are ambiguous and hard.

- 2D -> 3D
- Segmentation/Edges



So, how do we solve these problems?



- Magic of data !
- Use data to learn better likelihoods: how things look like.
- Use data to learn priors of what is more likely than others.

But how much data do we need?

The extremes of learning

Extrapolation problem
Generalization
Transfer learning

Interpolation problem
Correspondence
Finding the differences



So how much data does humans use?

What's the Capacity of Visual Long Term Memory?

What we know...

Standing (1973)

10,000 images

83% Recognition

... people can remember thousands of images

High Fidelity Visual Memory is possible
(Hollingworth 2004)

What we don't know...

... what people are remembering for each item?




According to Standing

"Basically, my recollection is that we just separated the pictures into **distinct thematic categories**: e.g. cars, animals, single-person, 2-people, plants, etc.) Only a few slides were selected which fell into each category, and they were visually distinct."



Dogs
Playing Cards

"Gist" Only



Sparse Details



Highly Detailed

Massive Memory I: Methods



Showed 14 observers 2500 **categorically unique objects**

1 at a time, 3 seconds each

800 ms blank between items

Study session lasted about 5.5 hours

Repeat Detection task to maintain focus

Followed by 300 2-alternative forced choice tests



how far can we push the fidelity of visual LTM representation ?

Same object category, different instance




how far can we push the fidelity of visual LTM representation ?


Same object, different states




Massive Memory I: Recognition Memory Results



Massive Memory I: Recognition Memory Results



Extrapolation of Repeat Detection Data



how much data does computer vision
researchers use?

10^0
images



1972

10^1
images



10^1
images





FIGURE 4. This figure provides a high quality reproduction of the six images discussed in the text. (a) and (b) were taken with a considerably modified Information International Incorporated Vidissector, and the rest were taken with a Telemation TMC-2100 vidicon camera attached to a Spatial Data Systems digitizer (Camera Eye 108). The full dynamic range from black to white is represented by 256 grey-levels. The images reproduced here were created by an Optronics P1500hPhotowriter from intensity arrays that measured 128 elements square. This size of intensity array corresponds to viewing a 1 in square at 5 ft with the human retina. The image of the period at the end of this sentence probably covers more than 40 retinal receptors. The reader should view the images from a distance of about 5 ft when assessing the performance of the programs.


Marr, 1976

10^{2-4}
images




The faces and cars scale

10^{2-4}
images



In 1996 DARPA released 14000 images, from over 1000 individuals.



The PASCAL Visual Object Classes

In 2007, the twenty object classes that have been selected are:

Person: person

Animal: bird, cat, cow, dog, horse, sheep

Vehicle: aeroplane, bicycle, boat, bus, car, motorbike, train


Indoor: bottle, chair, dining table, potted plant, sofa, tv/monitor



10^{2-4}
images



10^5
images



10^5
images

Caltech 101 and 256



Fei-Fei, Fergus, Perona, 2004



Griffin, Holub, Perona, 2007

Lotus Hill Research Institute image corpus





Figure 5: Two examples of the parse trees (cat and car) in the Lotus Hill Research Institute image corpus. From [87].

LabelMe

10^5
images





Tool went online July 1st, 2005
530,000 object annotations collected

Labelme.csail.mit.edu

B.C. Russell, A. Torralba, K.P. Murphy, W.T. Freeman, IJCV 2008

Quality of labeling



25%

50%

75%

Average labeling quality

Extreme labeling



The other extreme of extreme labeling

... things do not always look good...



Creative testing

LabelMe Name [suggestions](#) if you find any bugs or have any suggestions.

Scan in [help] Show me another image

Label as many objects and regions as you can in this image

There are 146302 labelled objects

Instructions [Get more help]

Use your mouse to click around the boundary of some objects in this image. You will then be asked to enter the name of the object (examples: car, window).

Labeling tools

[Erase polygons](#) [Zoom](#) [Image](#) [+](#) [-](#)


Polygons in this image

[Done](#)

Regions [Background](#) [Floor](#) [Wall](#) [Chair](#) [Table](#) [Couch](#) [Bed](#)




Scene and object biases



10^5
images



10⁶⁻⁷
images





Things start getting out of hand

Collecting big datasets

10^{6-7}
images

- **ESP game (CMU)**
Luis Von Ahn and Laura Dabbish 2004
- **LabelMe (MIT)**
Russell, Torralba, Freeman, 2005
- **StreetScenes (CBCL-MIT)**
Bileschi, Poggio, 2006
- **WhatWhere (Caltech)**
Perona et al, 2007
- **PASCAL challenge**
2006, 2007
- **Lotus Hill Institute**
Song-Chun Zhu et al, 2007
- **80 million images**
Torralba, Fergus, Freeman, 2007





80.000.000 images

10^{6-7}


75.000 non-abstract nouns from WordNet

7 Online image search engines^{images}




And after 1 year downloading images

Google: 80 million images



- An ontology of images based on WordNet
- ImageNet currently has
 - 22,000+ categories of visual concepts
 - 15 million human-cleaned images (~700im/categ)
 - 1/3+ is released online @ www.image-net.org



Mechanical Turk is a marketplace for work.

We give businesses and developers access to an on-demand, scalable workforce.

Workers select from thousands of tasks and work whenever it's convenient.

216,070 HITs available. [View them now.](#)

Make Money by working on HITs

HITs - Human Intelligence Tasks - are individual tasks that you work on. [Find HITs now.](#)

As a Mechanical Turk Worker you:

- Can work from home
- Choose your own work hours
- Get paid for doing good work

Find an
interesting task

Work

Earn
money



[Find HITs Now](#)

or [learn more about being a Worker](#).

Get Results from Mechanical Turk Workers

Ask workers to complete HITs - Human Intelligence Tasks - and get results using Mechanical Turk. [Register Now](#)

As a Mechanical Turk Requester you:

- Have access to a global, on-demand, 24 x 7 workforce
- Get thousands of HITs completed in minutes
- Pay only when you're satisfied with the results

Fund your
account



Load your
tasks



Get
results



[Get Started](#)

Labeling for money



Bryan C Russell | Account Settings | Sign Out | Help

Your Account

HITs

Qualifications

56,035 HITs
available now

All HITs | HITs Available To You | HITs Assigned To You

Search for HITs containing

that pay at least \$ 0.00 for which you are qualified

Timer: 00:00:13 of 60 minutes

Finished with this HIT? Let someone else do it?

Submit HIT

Reopen HIT

Automatically accept the next HIT

Total Earned: \$0.01

Total HITs Submitted: 12

LabelMe: Label objects in the image

Requester: Bryan C Russell

Qualifications Required: None


Rewards: \$0.01 per HIT

HITs Available: 250

Durations: 60 minutes


Please label as many objects as you want in this image. Scroll down to see the entire image.

Submit HIT



Alexander Sorokin, David Forsyth, "Utility data annotation with Amazon Mechanical Turk", First IEEE Workshop on Internet Vision at CVPR 08.

10^{6-7}
images



10⁸⁻¹¹
images



Datasets in perspective

Number of images on my hard drive: 10^5

Number of images seen during my first 10 years: 10^8
($3 \text{ images/second} * 60 * 60 * 16 * 365 * 10 = 630720000$)

Number of images seen by all humanity: 10^{20}

$106,456,367,669 \text{ humans}^1 * 100 \text{ years} * 3 \text{ images/second} * 60 * 60 * 16 * 365 =$
1 from <http://www.prb.org/Articles/2002/HowManyPeopleHaveEverLivedonEarth.aspx>

Number of all 32x32 images: 10^{7373}

$256^{32*32*3} \sim 10^{7373}$




When do we need big data?

Two Kinds of Things in the World



Navier-Stokes Equation

$$\frac{\partial \mathbf{u}}{\partial t} = -(\mathbf{u} \cdot \nabla) \mathbf{u} + v \nabla^2 \mathbf{u} - \frac{1}{d} \nabla p + \mathbf{f}$$



+ weather
+ location
+ ...

Lots of data available

Signed in as [swotjerel](#) [Edit Profile](#)

flickr now you can!

[Home](#) [You](#) [Organize & Create](#) [Contacts](#) [Groups](#) [Explore](#)

Search [Photos](#) [Groups](#) [People](#)

[Full Text](#) | [Tags Only](#) [Advanced Search](#)

Sort: [Relevant](#) | [Recent](#) | [Interesting](#)

View: [Small](#) | [Medium](#) | [Detail](#) | [Slideshow](#)

 From Odalagh	 From coat123	 From Martin LaBar...	 From floomeschb...	 From *Iris-hued*	 From Fareed...	 From Gary...
 From ExeDaro	 From Mark_Walor...	 From hogsvilleBr...	 From Photo by...	 From martyspants...	 From dnixel...	 From Lord V
 From ConnellPK....	 From e_shark	 From RasA	 From mark...	 From calbergem...	 From SixRevisions	 From GrungeTextur...
 From ConnellPK....	 From e_shark	 From RasA	 From calbergem...	 From mark...	 From SixRevisions	 From SixRevisions

“Unreasonable Effectiveness of Data”

[Halevy, Norvig, Pereira 2009]

- Parts of our world can be explained by elegant mathematics:
 - physics, chemistry, astronomy, etc.
- But much cannot:
 - psychology, genetics, economics,... visual understanding?
- Enter: The Magic of Data
 - Great advances in several fields:
 - e.g. speech recognition, machine translation, Google

Unreasonable Effectiveness of Data

Simple Algorithms (Dumb) + Lot of Data
are better than Complicated algorithms


Example: Machine Translation

Example: Texture Generation

Machine Translation

Step 1: Source Sentence Chunking


- Segment source sentence into overlapping n-grams via sliding window
- Typical n-gram length 4 to 9 terms
- Each term is a word or a known phrase
- Any sentence length



Step 2: Dictionary Lookup


- Using bilingual dictionary, list all possible target translations for each source word or phrase

Source Word-String




Step 3: Search Target Text

- Using the Flooding Set, search target text for word-strings containing one word from
 - each group




- Find maximum number of words from Flooding Set in minimum length word-string
 - *Words or phrases can be in any order*
 - *Ignore function words in initial step (T5 is a function word in this example)*


Step 3: Search Target Text (Example)



Step 3: Search Target Text (Example)



Step 3: Search Target Text (Example)




Reintroduce function words after initial match (e.g. T5)

Slide by Jaime Carbonell

Step 4: Score Word-String Candidates

- Scoring of candidates based on:
 - Proximity (minimize extraneous words in target n-gram \approx precision)
 - Number of word matches (maximize coverage \approx recall)
 - Regular words given more weight than function words
 - Combine results (e.g., optimize F_1 or p-norm or ...)

Target Word-String Candidates



Step 5: Select Candidates Using Overlap (Propagate context over entire sentence)

Word-String 1
Candidates

T(x1)	T2-d T3-c T(x2)	T4-b
T(x1)	T3-c T2-b T4-e	
T(x2)	T4-a T6-b T(x3)	T2-c

Word-String 2
Candidates

T3-b T(x3)	T2-d T(x5)	T(x6)	T6-c
T4-a T6-b T(x3)	T2-c T3-a		
T3-c T2-b T4-e T5-a T6-a			


Word-String 3
Candidates

T2-b T4-e T5-a T6-a T(x8)
T6-b T(x11) T2-c T3-a T(x9)
T6-b T(x3) T2-c T3-a T(x8)


Step 5: Select Candidates Using Overlap

Best translations selected via maximal overlap

Alternative 1



Alternative 2




A (Simple) Real Example of Overlap

Flooding → N-gram fidelity

Overlap → Long range fidelity

N-grams
generated
from
Flooding



a United States soldier

United States soldier died

soldier died and two others

died and two others were injured

two others were injured Monday

N-grams connected via
Overlap

a United States soldier died and two others were injured Monday

Texture Synthesis

Texture Synthesis



Classical Texture Synthesis

Synthesis

Novel texture


Parametric
Texture
Model

Analysis

Sample texture

This is hard!

Throwing away too much too soon?



input texture



synthesized texture

Non-parametric Approach

Synthesis


Novel texture





Analysis

Sample texture

[Efros & Leung, '99, Efros & Freeman '01]




Texture Growing




So, how do we use big data?

Two ways to use Lots of Data




Brute Force Vision: Find that needle in the haystack and disregard the rest (a.k.a. kNN)




See what different subsets of data think of you

kNN matching is great...


- because we live in a (mostly) boring world!




Lots Of Images



Lots Of Images



Lots Of Images




Automatic Colorization Result

Grayscale input High resolution



Colorization of input using average



im2gps

Instead of using objects labels, the web provides other kinds of metadata associate to large collections of images




Figure 2. The distribution of photos in our database. Photo locations are cyan. Density is overlaid with the jet colormap (log scale).

20 million geotagged and geographic text-labeled images

im2gps

Hays & Efros. CVPR 2008




Figure 5. *Geolocation performance across features.* Percentage of test cases geolocated to within 200km for each feature. We compare geolocation by 1-NN vs. largest mean-shift mode.




Image completion



Original Image

Input

Criminisi et al.

MS Smart Erase

Instead, generate proposals using millions of images



Input

16 nearest neighbors
(gist+color matching)

output

Hays, Efros, 2007

With a good image similarity
and a lot of data...

Input image




22,000 LabelMe scenes

Nearest neighbors








Hays, Efros, Siggraph 2006
Russell, Liu, Torralba, Fergus, Freeman. NIPS 2007

With a good image similarity and a lot of data...






With a good image similarity and a lot of data...







Russell, Liu, Torralba, Fergus, Freeman. NIPS 2007

With a good image similarity and a lot of data...





With a good image similarity and a lot of data...



Russell, Liu, Torralba, Fergus, Freeman. NIPS 2007


Outputs



While many scenes are boring...




Some scenes are unique



Slide by Antonio Torralba

Dealing with sparse data (rare scenes)

- better similarity



Medici Fountain, Paris



Search by image



Drop image here

 Move

[Watch a short video to learn more.](#)



medici_summer.jpg [X](#)

luxembourg gardens



Search

About 2 results (0.29 seconds)

Everything

Images

Maps

Videos

News

Shopping

More



Image size:
1024 × 829

No other sizes of this image found.

Visually similar





Medici Fountain, Paris (winter)



medici_winter.png



luxembourg gardens



Search

About 2 results (0.29 seconds)

Everything

Images

Maps

Videos

News

Shopping

More



Image size.

713 × 600

No other sizes of this image found.

Visually similar







painting.png

describe image here



Search

About 2 results (0.29 seconds)

[Everything](#)[Images](#)[Maps](#)[Videos](#)[News](#)[Shopping](#)[More](#)

Image size:
319 × 482

No other sizes of this image found.

Visually similar







medici_sketch.bmp

describe image here



Search




About 2 results (0.29 seconds)

[Everything](#)[Images](#)[Maps](#)[Videos](#)[News](#)[Shopping](#)[More](#)


Image size
413 × 182


No other sizes of this image found.

Visually similar




OUR GOAL






Input Query




Top Matches

Input Query



Top Matches


Input Query



Top Matches

IMPORTANT PARTS?

Input Query





Important Parts



Top Matches

Input Query





“Data-driven Uniqueness”



Search using Images

Input Query



Top Matches

100

Search using Sketches



Search using Paintings



Input Painting



Top Matches

Search using Paintings



Input Painting



Top Matches

Dealing with sparse data (rare scenes)

- better similarity
- better alignment
 - e.g. reduce resolution, sifting, warping, etc.

Matching scenes

Two images taken from the same scene category, but different instances

- Contain different objects with different scales, perspectives and spatial location



Image representation










Image gradients




Keypoint descriptor





Scene parsing results





Dealing with sparse data (rare scenes)

- better similarity
- better alignment
 - e.g. reduce resolution, sifting, warping, etc.
- Use sub-images (primitives) to match
 - Allows matching from multiple images



Predicting Surface Normals




Matching Parts



Matching Parts



Matching Parts



Dealing with sparse data (rare scenes)

- better similarity
- better alignment
- Use sub-images (primitives) to match
- Understand the simple stuff first
 - e.g. tracking via recognition, background subtraction, “object pop-out”, etc.

Recognize when it's easy!

People take on a variety of **poses**, aspects, scales



self-occlusion



rare pose



motion blur



non-distinctive pose




too small




just right
detect this²⁷


Guess structure




Guess structure




Subtracting away structure



Structure




Objects




Wall appearance modeling

Dealing with sparse data (rare scenes)

- better similarity
- better alignment
 - e.g. reduce resolution, sifting, warping, etc.
- segment into chunks
 - e.g. segmentation for recognition approaches
- get rid of simple stuff first
 - e.g. background subtraction, “object pop-out”, etc.
- Moving away from kNN methodology...
- use data to make connections
 - e.g. The Memex, manifold learning, data association, subpopulation means, etc.



Memex - Knowledge Graph




Manifolds

- Images are high dimensional: A 64x64 image is 4096 dimensional vector.
- But the possible images are much less!
- Is there a subspace where the set of images lie?

manifolds in vision

appearance variation




manifolds in vision




Slide by Dave Thompson

reasonable distance metrics




reasonable distance metrics



linear interpolation


reasonable distance metrics




manifold interpolation

Slide by Dave Thompson

reasonable distance metrics




reasonable distance metrics




Some observations about data collection

Object distributions




LabelMe




IMAGENET

SUN database




The first 9 objects account for 50% of all training examples
17 classes with more than 300 examples
109 classes with less than 50 examples

Rare objects are similar to frequent objects



Rare objects are similar to frequent objects



Some bias comes from the way the data is collected

mug

Search

Refine search results

About 10,100,000 results (0.00 seconds)

Advanced search

\$8.99 Logo Coffee Mugs

www.DiscountMugs.com Lead Free & Dishwasher Safe. Save 40-50%. No Catch. Factory Direct!

Custom Mugs On Sale

www.Vistaprint.com Order Now & Save 50% On Custom Mugs. No Minimums. Upload Photos & Logos.

Promotional Mugs from 69¢

www.4imprint.com/Mugs Huge Selection of Stylish Colors. Buy 72 Mugs @ \$1.95 ea. 24hr Service!

Related searches: white mug coffee mug mug root beer mug shot



Representational
500 × 429 - 81k - jpg
egegneyes.org
Find similar Images



Ceramic Happy Face
300 × 300 - 77k - jpg
laruse.com
Find similar Images



Here I go then, trying
800 × 800 - 35k - jpg
beeper.wordpress.com
Find similar Images



The Chalk Mug
304 × 314 - 17k - jpg
coolest-gadgets.com
Find similar Images



mug



Bring your own
500 × 451 - 15k - jpg
www.yourmug.co.uk
Find similar Images



ceramic mug
980 × 1024 - 30k - jpg
diytrade.com



Dual Purpose Drinking
480 × 428 - 18k - jpg
freshome.com
Find similar Images



This coffee mug.
300 × 300 - 22k - jpg
gizmodo.com
Find similar Images



Back to Ceramic
400 × 400 - 8k - jpg
hashpromotions.com.au
Find similar Images



Coffee Mug as a
308 × 301 - 10k - jpg
duoshui.wordpress.com
Find similar Images




SASS Life Member
300 × 302 - 8k - jpg
sassanet.com



personalized coffee
400 × 343 - 15k - jpg
wallyou.com
Find similar Images



We like our mugs
290 × 290 - 8k - jpg
kthenancynotions.com
Find similar Images



TEAPOT

[\\$9.95 Logo Coffee Mugs](#)[www.DiscounMugs.com](#) Lead Free & Dishwasher Safe. Save 40-50%. No Catch. Factory Direct![Custom Mugs On Sale](#)[www.Vistaprint.com](#) Clear Now & Save 50% On Custom Mugs. No Minimums. Upload Photo & Logos.[Promotional Mugs from \\$9.95](#)[www.4imp.com/Mugs](#) Huge Selection of Styles Colors- Buy 72 Mugs @ \$1.35 ea-24hr Service

Google mugs

Related searches: white mug coffee mug mug root beer mug shot



Representational
500 × 429 - 81k - jpg
eegereyes.org
Find similar images



Ceramic Happy Face
300 × 300 - 77k - jpg
larasee.com
Find similar images



Here I go then, trying
800 × 800 - 35k - jpg
deeper.wordpress.com
Find similar images



The Chalk Mug
304 × 314 - 17k - jpg
coolest-gadgets.com
Find similar images



mug
300 × 279 - 54k - jpg
reymerwatch.org



Bring your own
500 × 451 - 15k - jpg
cookswomnuked.co
Find similar images



ceramic mug
800 × 1024 - 30k - jpg
diytrade.com



Dual Purpose Drinking
400 × 420 - 18k - jpg
inachoice.com
Find similar images



This coffee mug
300 × 300 - 22k - jpg
giamodo.com
Find similar images



Back to Ceramic
400 × 400 - 8k - jpg
meanpromotion.com.au
Find similar images




Coffee Mug ss s
303 × 301 - 10k - jpg
diythewordpress.com
Find similar images



SASS Life Member
300 × 302 - 8k - jpg
sasseanet.com
Find similar images




personalized coffee
400 × 343 - 15k - jpg
wallyou.com
Find similar images



Mugs from LabelMe

“Name That Dataset!” game



Caltech 101

Caltech 256

MSRC

UIUC cars

Tiny Images

Corel

PASCAL 2007

LabelMe

COIL-100

ImageNet

15 Scenes

SUN'09


SVM plays “*Name that dataset!*”

SVM plays “Name that dataset!”

	UIUC	LabelMe	PASCAL07	MSRC	SUN09	15 Scenes	Corel	Caltech101	Caltech256	Tiny	ImageNet	COIL-100
UIUC	0	29	8	21	3	10	2	17	8	1	2	0
LabelMe	0	54		7	8	6		2	2	4	6	0
PASCAL 2007	0	10	29	10	10		7		7	7	11	1
MSRC	0	3	7	60		3	1		2		7	0
SUN09	0	14	9	9	24	17	11	4	3	4	1	0
15 Scenes	0	8	3		13	51	11	2	2	4	2	0
Corel	1	2	8		8	11	35	10	7	7	9	0
Caltech101	1	2	9	9	2	4	7	38	14	7	6	1
Caltech256	1	2	8			6	10	18	20	11	12	1
Tiny	1	2	8	8	5	6	11	12	13	24	12	1
ImageNet	1	3	11	9	8	4	11	8	12	13	21	1
COIL-100	0	0	0	0	0	0	0	0	0	0	0	99

- 12 1-vs-all classifiers
- Standard full-image features
- 39% performance (chance is 8%)

SVM plays “Name that dataset!”



Datasets have different goals...

- Some are object-centric (e.g. Caltech, ImageNet)
- Otherwise are scene-centric (e.g. LabelMe, SUN'09)
- What about playing “*name that dataset*” on bounding boxes?

Cross-Dataset Generalization

MSRC



Classifier trained on MSRC cars

Mixing datasets

Test on PASCAL

